

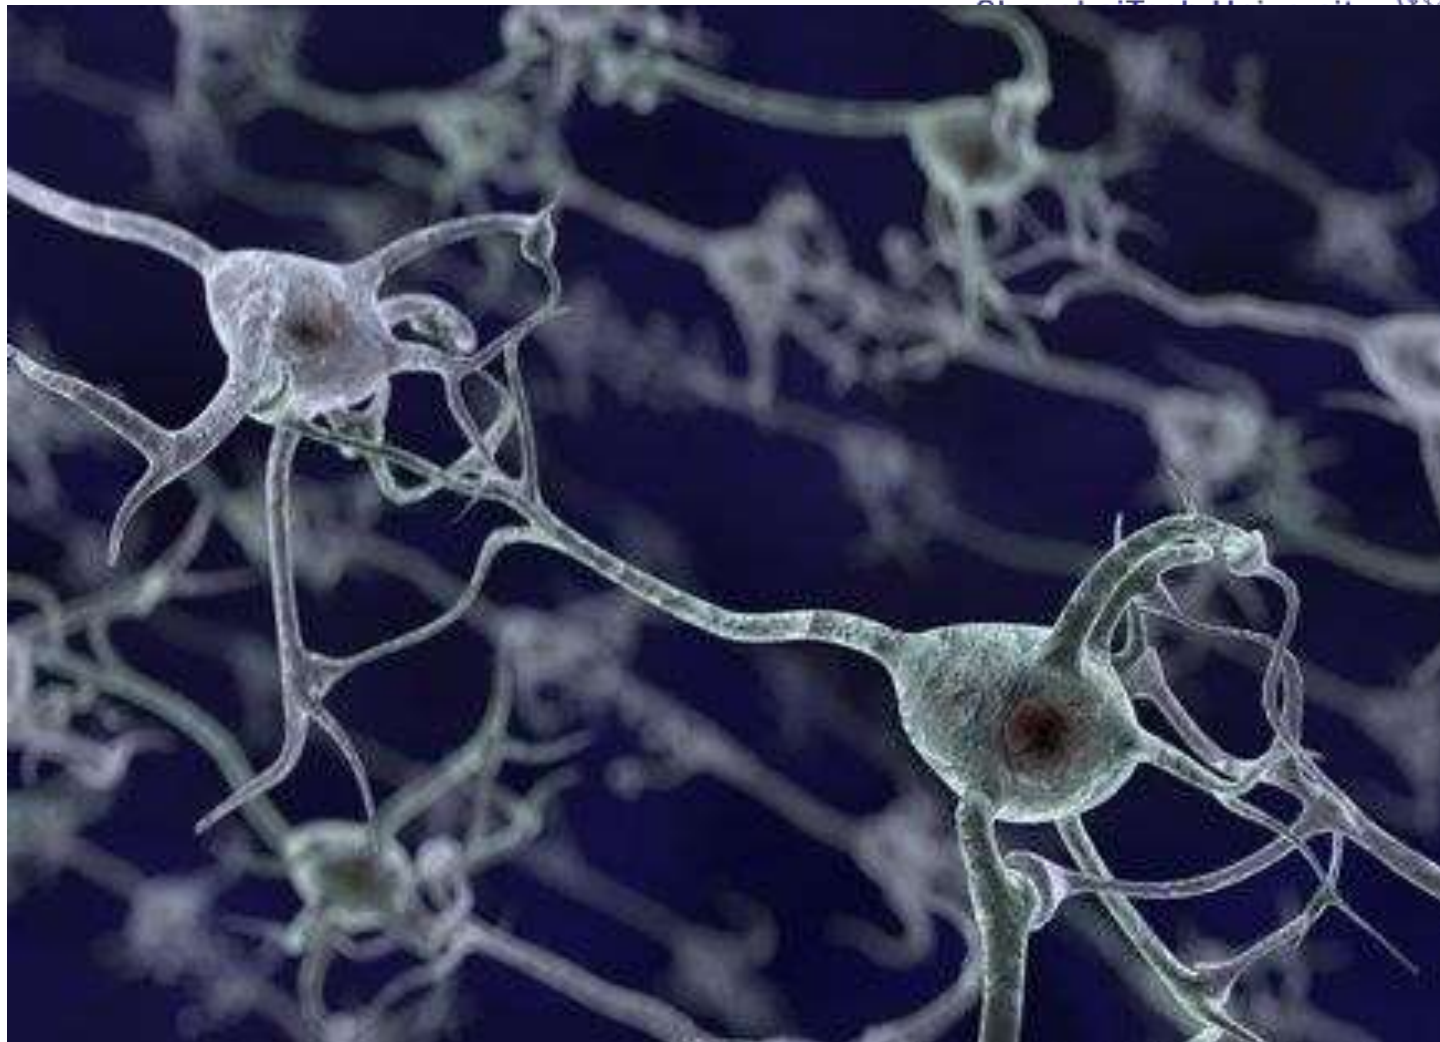


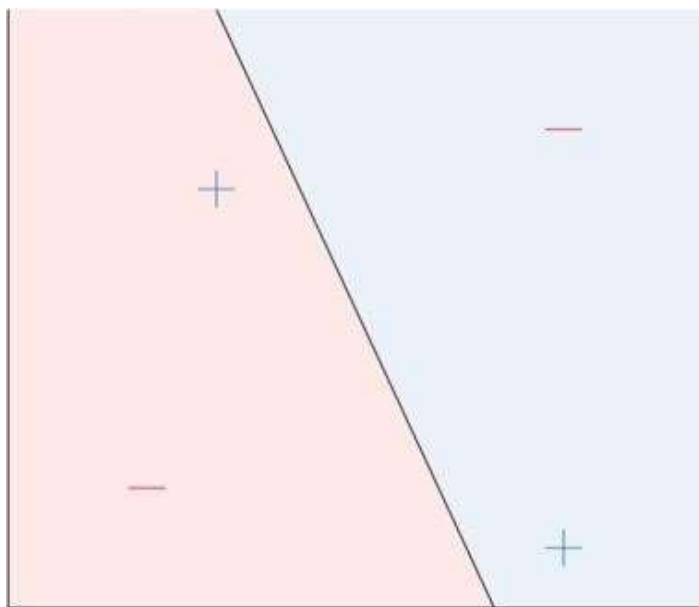
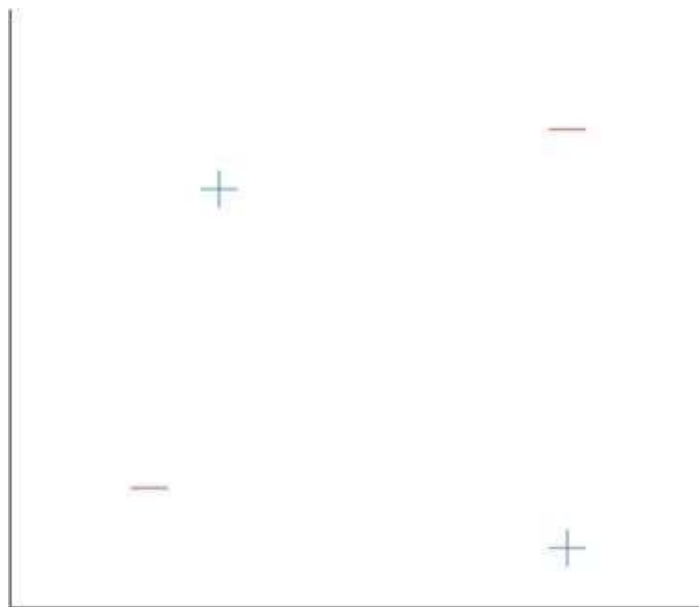
# CS182: Introduction to Machine Learning – Neural Networks

Yujiao Shi  
SIST, ShanghaiTech  
Spring, 2025

Slides Courtesy of Matt Gormley & Henry Chai, CMU

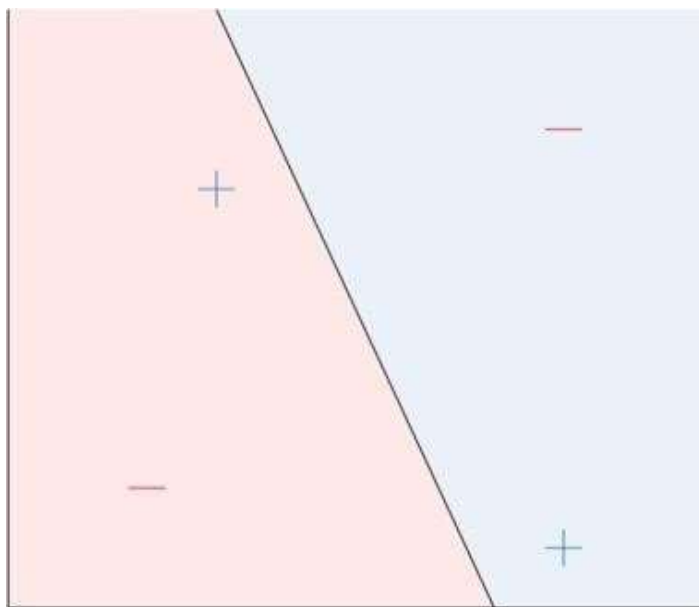
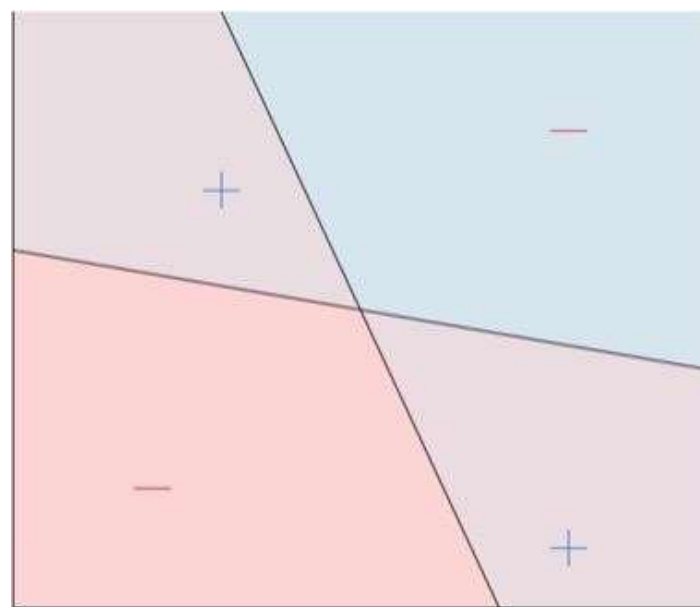
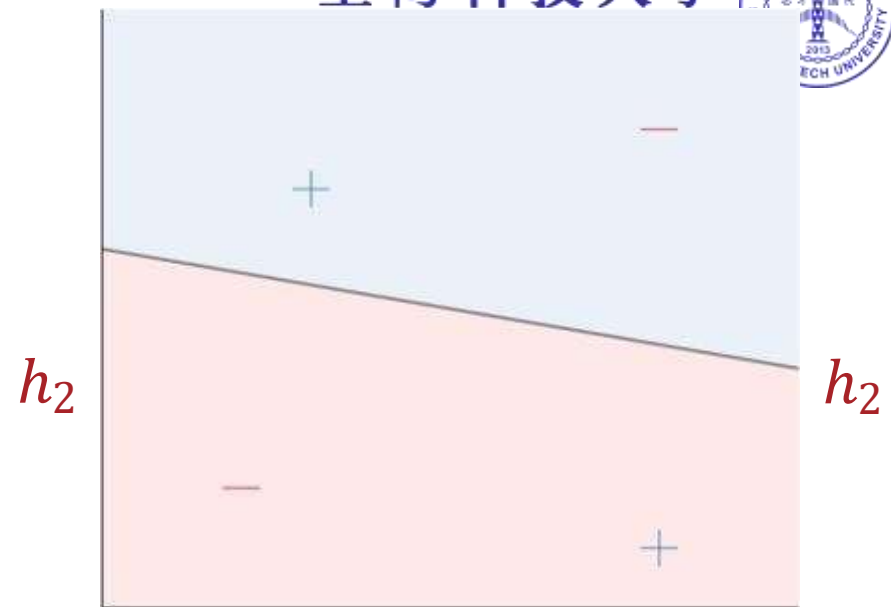
# Biological Neural Network



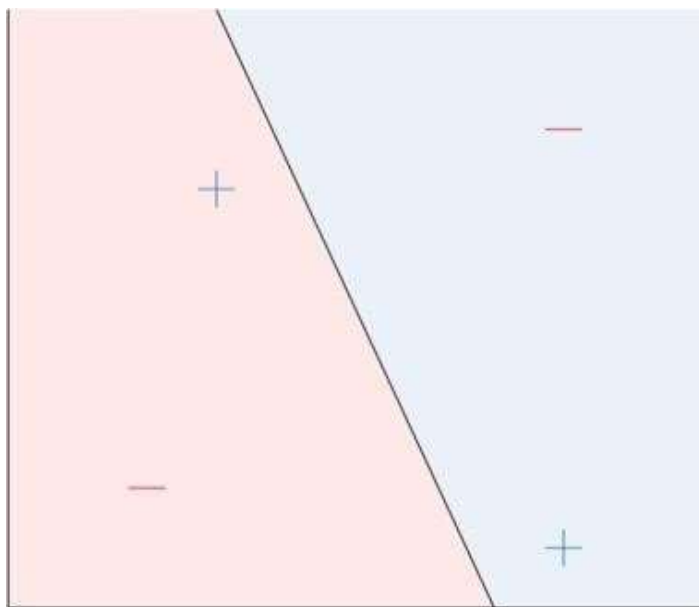
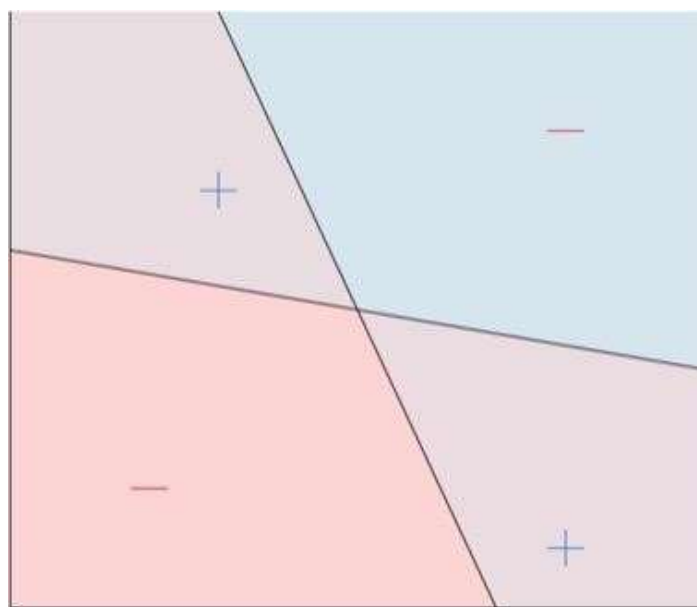
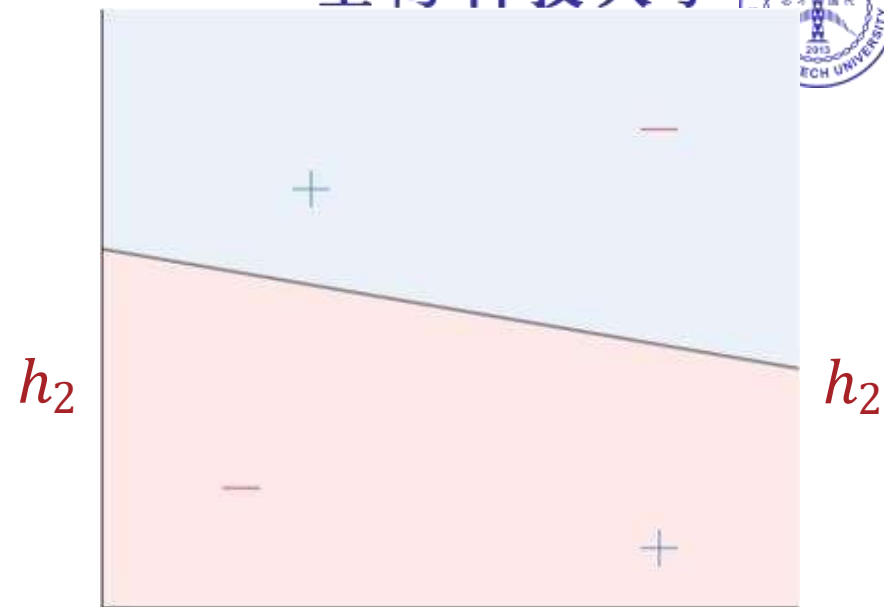
 $h_1$  $h_2$ 

# Perceptrons

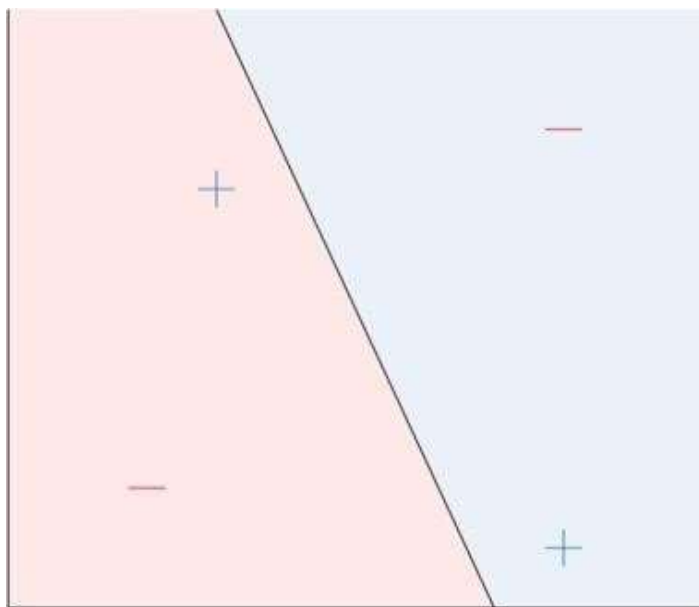
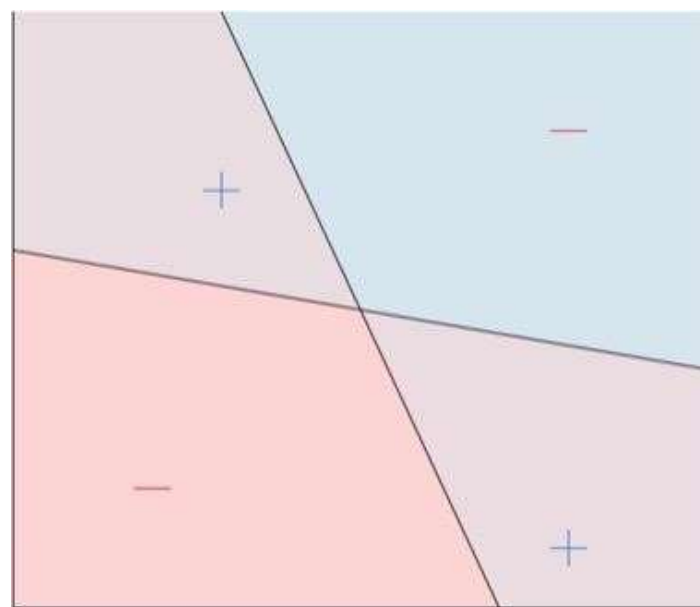
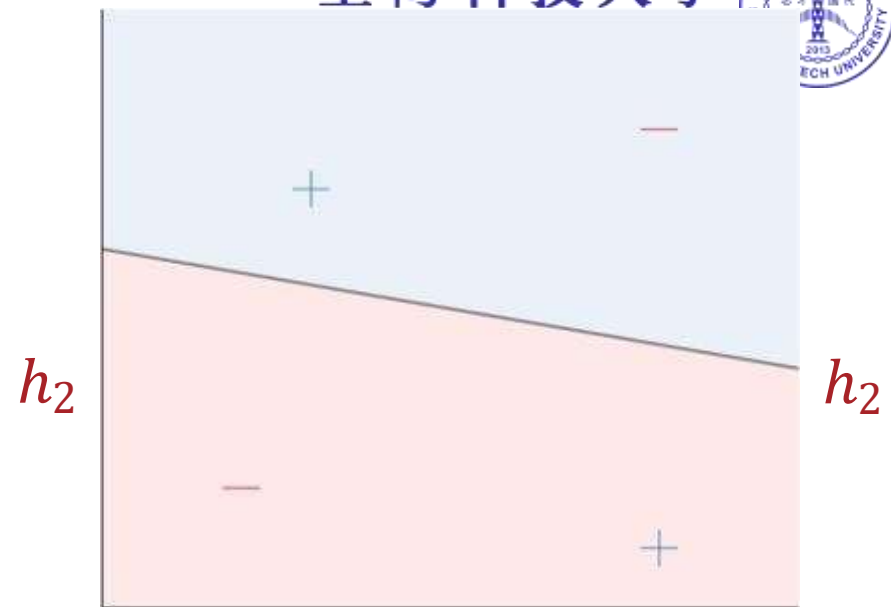
- Linear model for classification
- $h(\mathbf{x}) = \text{sign}(\mathbf{w}^T \mathbf{x})$
- Predictions are  $+1$  or  $-1$

 $h_1$  $h_1$  $h_2$ 

# Combining Perceptrons

 $h_1$  $h_1$  $h_2$ 

$$h(\mathbf{x}) = \begin{cases} +1 & \text{if } (h_1(\mathbf{x}) = +1 \text{ and } h_2(\mathbf{x}) = -1) \text{ or } (h_1(\mathbf{x}) = -1 \text{ and } h_2(\mathbf{x}) = +1) \\ -1 & \text{otherwise} \end{cases}$$

 $h_1$  $h_1$  $h_2$ 

$$h(x) = OR \left( AND(h_1(x), \neg h_2(x)), AND(\neg h_1(x), h_2(x)) \right)$$

# Boolean Algebra



- Boolean variables are either  $+1$  (“true”) or  $-1$  (“false”)
- Basic Boolean operations:
  - Negation:  $\neg z = -1 * z$
  - And:  $AND(z_1, z_2) = \begin{cases} +1 & \text{if both } z_1 \text{ and } z_2 \text{ equal } +1 \\ -1 & \text{otherwise} \end{cases}$
  - Or:  $OR(z_1, z_2) = \begin{cases} +1 & \text{if either } z_1 \text{ or } z_2 \text{ equals } +1 \\ -1 & \text{otherwise} \end{cases}$

# Boolean Algebra



- Boolean variables are either  $+1$  ("true") or  $-1$  ("false")
- Basic Boolean operations
  - Negation:  $\neg z = -1 * z$
  - And:  $AND(z_1, z_2) = \text{sign}(z_1 + z_2 - 1.5)$
  - Or:  $OR(z_1, z_2) = \text{sign}(z_1 + z_2 + 1.5)$



# Boolean Algebra



- Boolean variables are either  $+1$  ("true") or  $-1$  ("false")
- Basic Boolean operations
  - Negation:  $\neg z = -1 * z$

- And:  $AND(z_1, z_2) = \text{sign} \left( [-1.5, 1, 1] \begin{bmatrix} 1 \\ z_1 \\ z_2 \end{bmatrix} \right)$

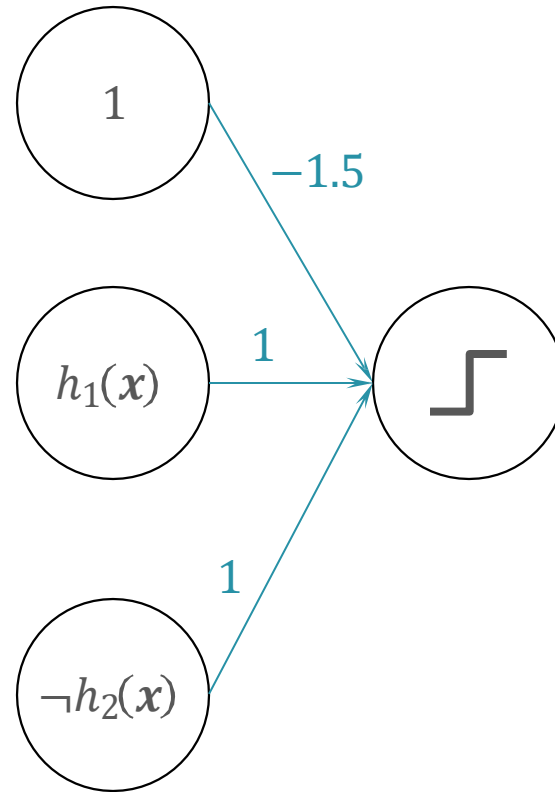
- Or:  $OR(z_1, z_2) = \text{sign} \left( [1.5, 1, 1] \begin{bmatrix} 1 \\ z_1 \\ z_2 \end{bmatrix} \right)$

## Building a Network

$$h(\mathbf{x}) = OR\left(AND(h_1(\mathbf{x}), \neg h_2(\mathbf{x})), AND(\neg h_1(\mathbf{x}), h_2(\mathbf{x}))\right)$$

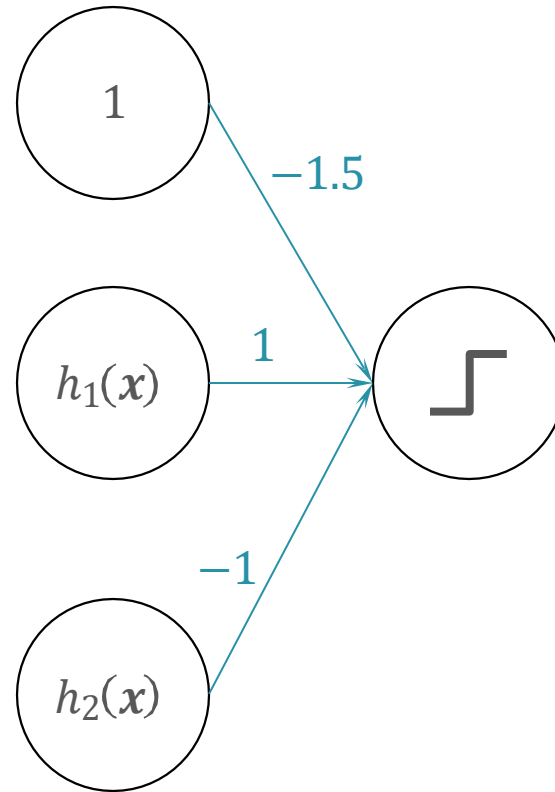
## Building a Network

$$h(x) = OR(AND(h_1(x), \neg h_2(x)), AND(\neg h_1(x), h_2(x)))$$



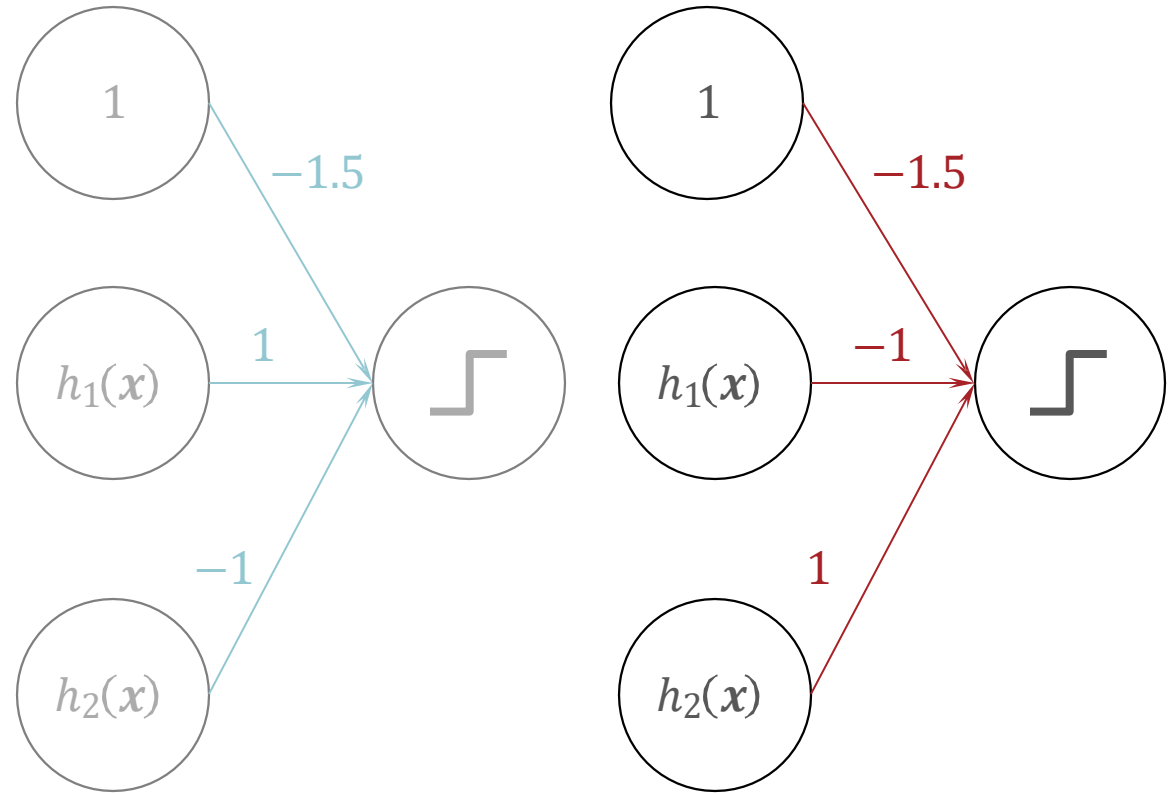
## Building a Network

$$h(x) = OR(AND(h_1(x), \neg h_2(x)), AND(\neg h_1(x), h_2(x)))$$



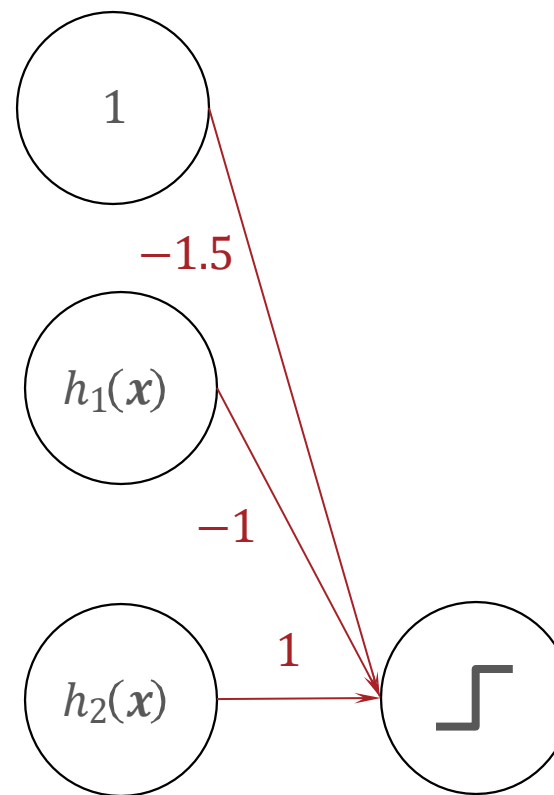
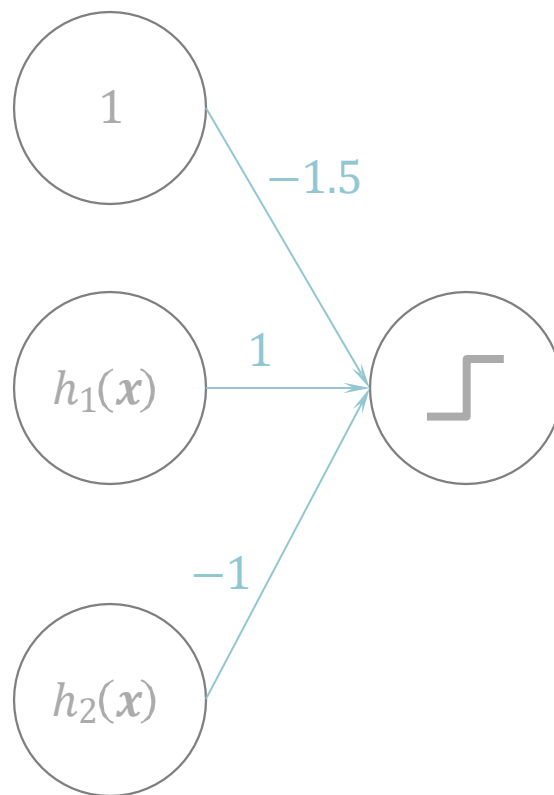
## Building a Network

$$h(x) = OR(AND(h_1(x), \neg h_2(x)), AND(\neg h_1(x), h_2(x)))$$



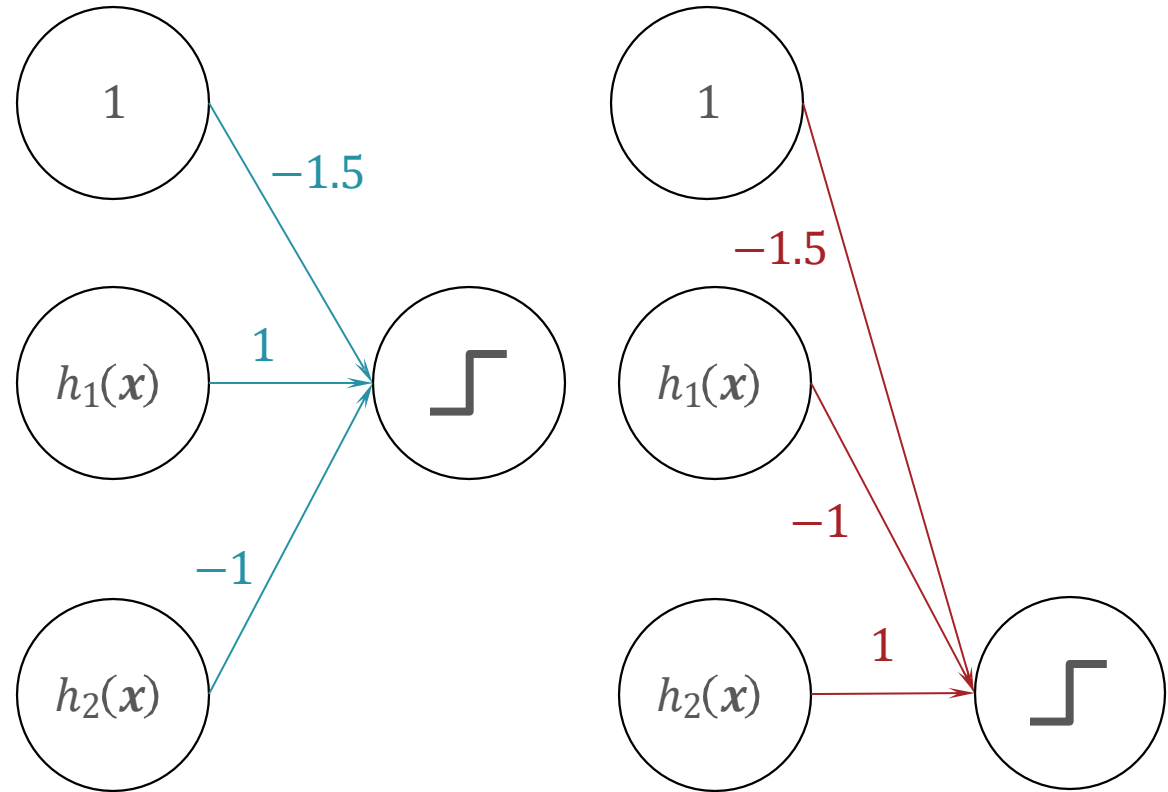
## Building a Network

$$h(x) = OR(AND(h_1(x), \neg h_2(x)), AND(\neg h_1(x), h_2(x)))$$



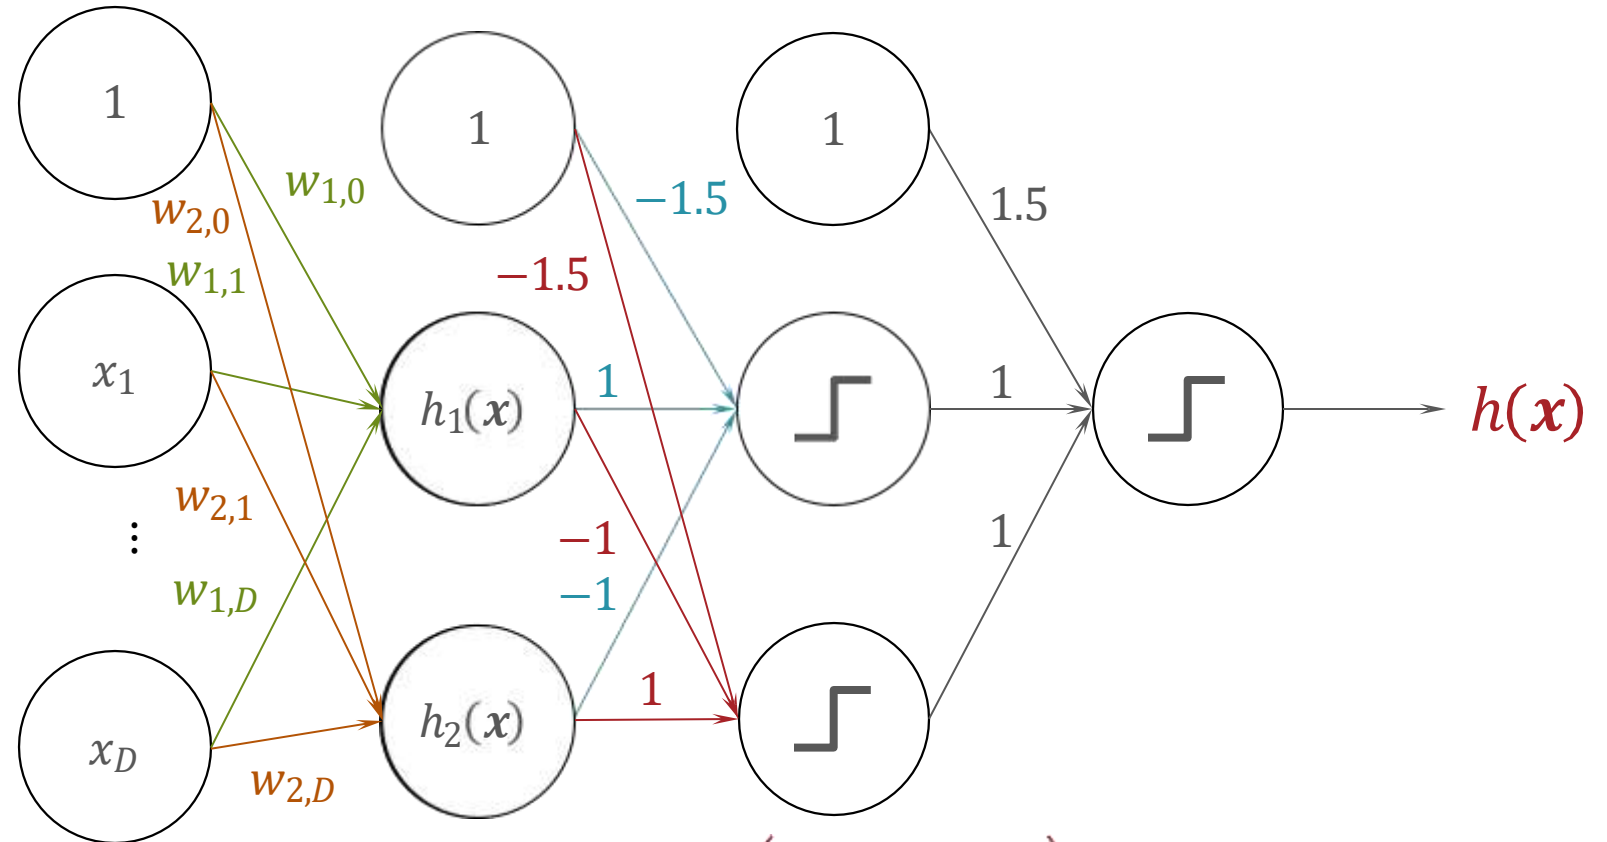
## Building a Network

$$h(x) = OR(AND(h_1(x), \neg h_2(x)), AND(\neg h_1(x), h_2(x)))$$



# Building a Network

$$h(\mathbf{x}) = OR(AND(h_1(\mathbf{x}), \neg h_2(\mathbf{x})), AND(\neg h_1(\mathbf{x}), h_2(\mathbf{x})))$$

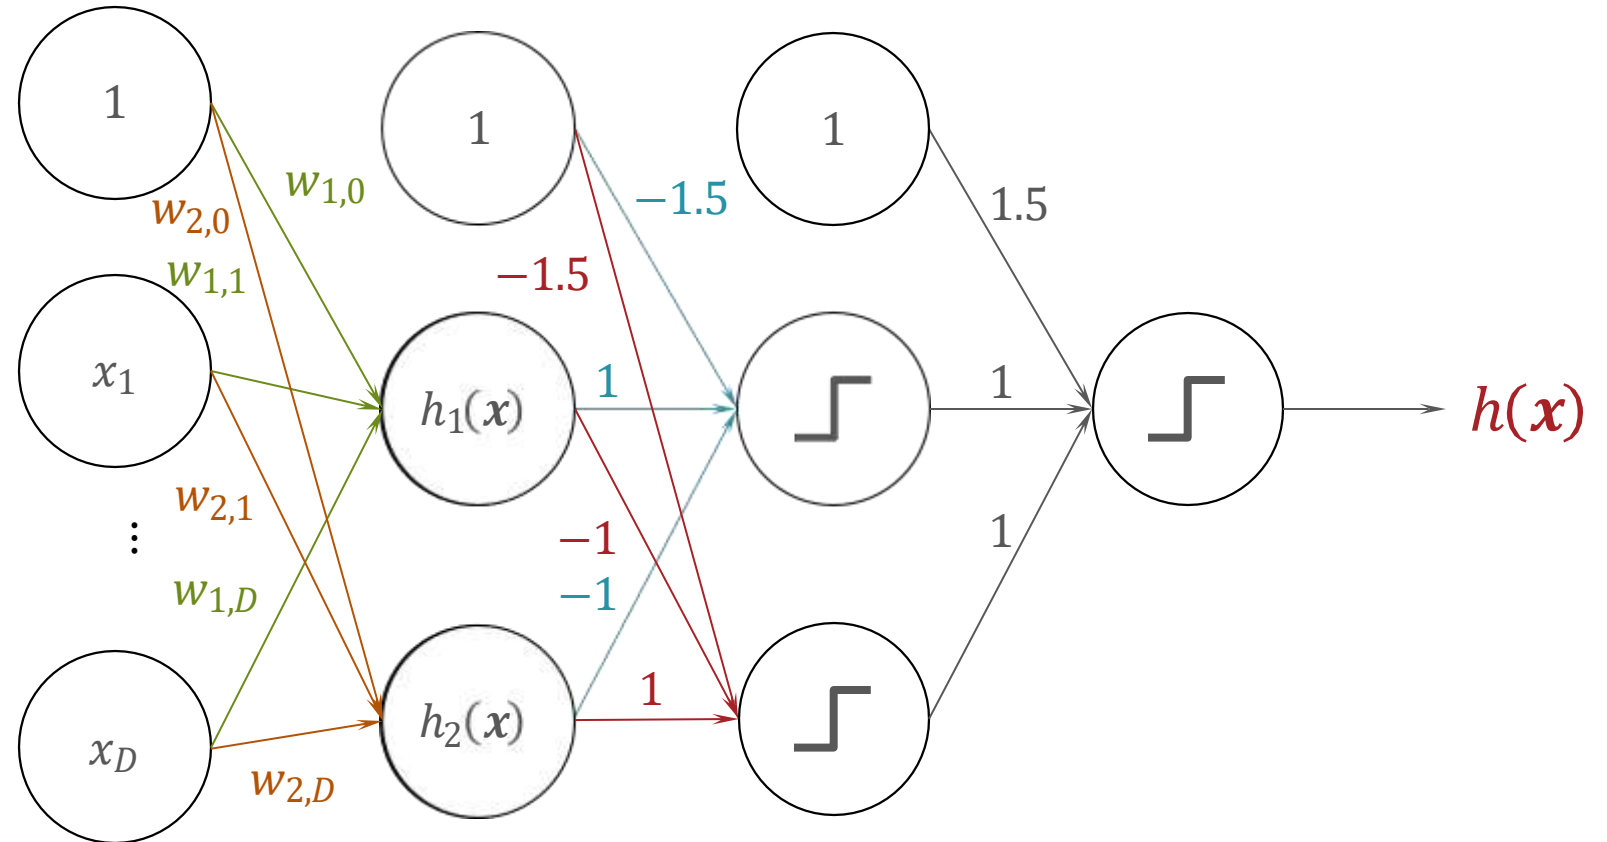


$$h_i(\mathbf{x}) = \text{sign}(\mathbf{w}_i^T \mathbf{x}) = \text{sign}\left(\sum_{d=0}^D w_{i,d} x_d\right)$$



# Building a Network

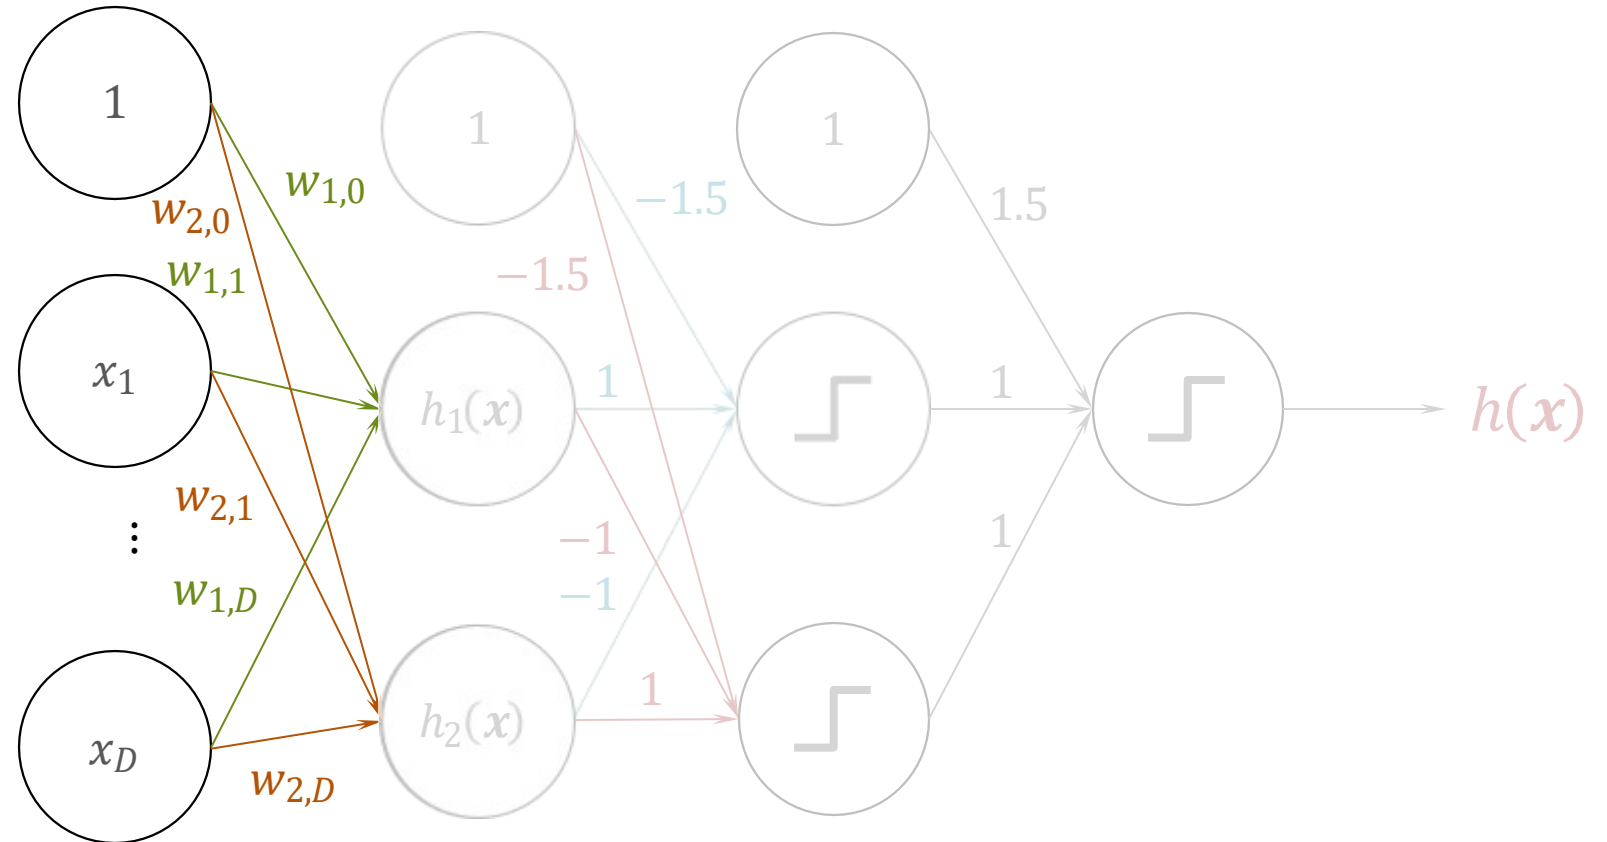
$$h(\mathbf{x}) = OR(AND(h_1(\mathbf{x}), \neg h_2(\mathbf{x})), AND(\neg h_1(\mathbf{x}), h_2(\mathbf{x})))$$



$$h(\mathbf{x}) = \text{sign}(\text{sign}(\text{sign}(\mathbf{w}_1^T \mathbf{x}) - \text{sign}(\mathbf{w}_2^T \mathbf{x}) - 1.5) + \text{sign}(-\text{sign}(\mathbf{w}_1^T \mathbf{x}) + \text{sign}(\mathbf{w}_2^T \mathbf{x}) - 1.5) + 1.5)$$

# Building a Network

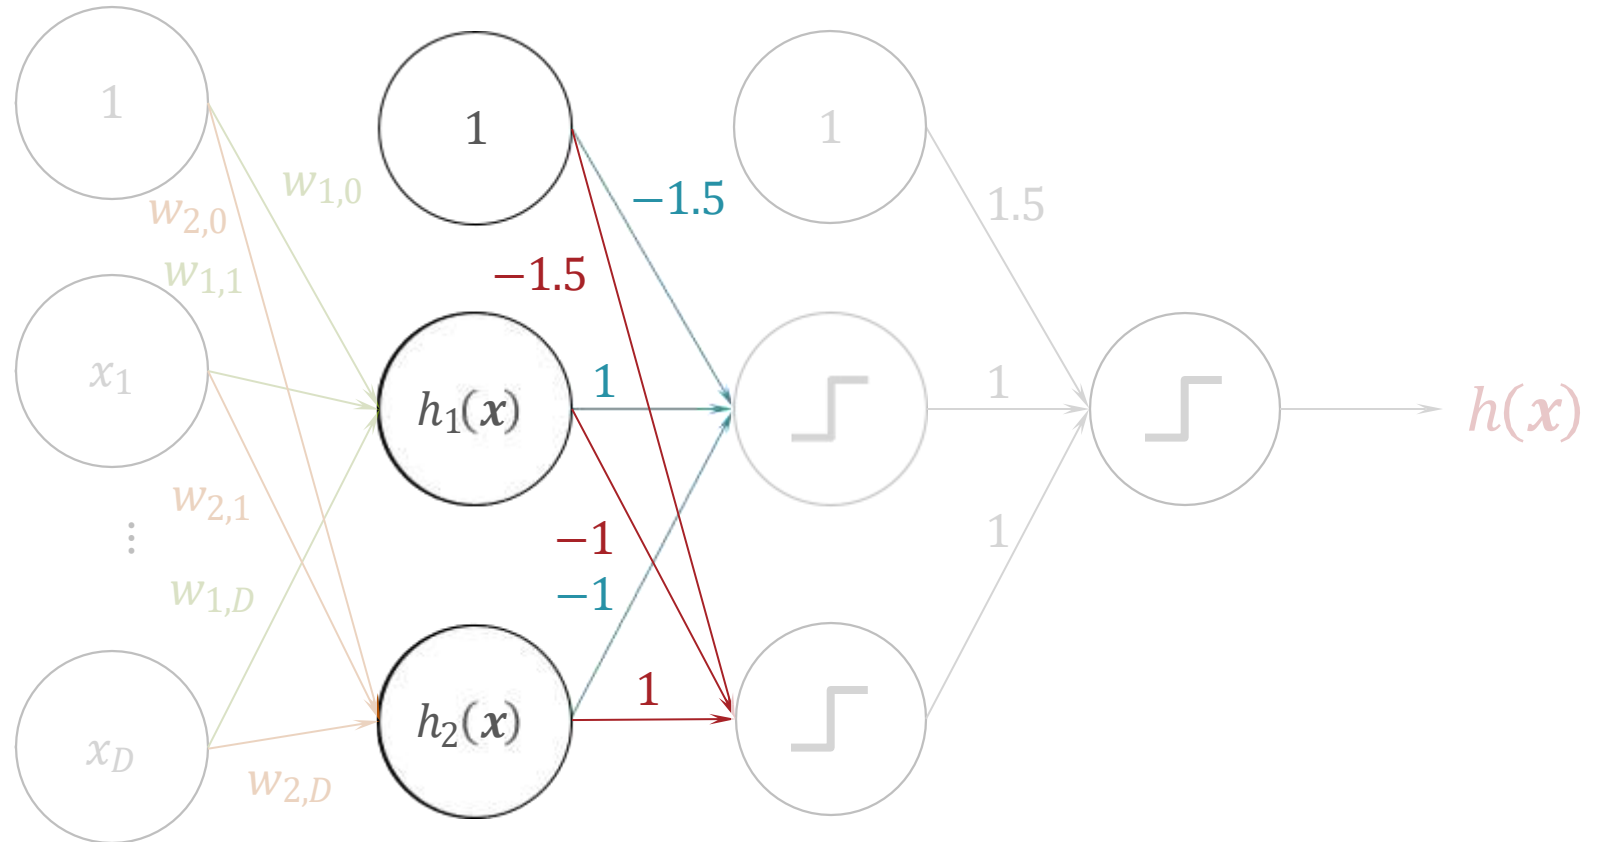
$$h(\mathbf{x}) = OR(AND(h_1(\mathbf{x}), \neg h_2(\mathbf{x})), AND(\neg h_1(\mathbf{x}), h_2(\mathbf{x})))$$



$$h(\mathbf{x}) = \text{sign}(\text{sign}(\text{sign}(\mathbf{w}_1^T \mathbf{x}) - \text{sign}(\mathbf{w}_2^T \mathbf{x}) - 1.5) + \text{sign}(-\text{sign}(\mathbf{w}_1^T \mathbf{x}) + \text{sign}(\mathbf{w}_2^T \mathbf{x}) - 1.5) + 1.5)$$

# Building a Network

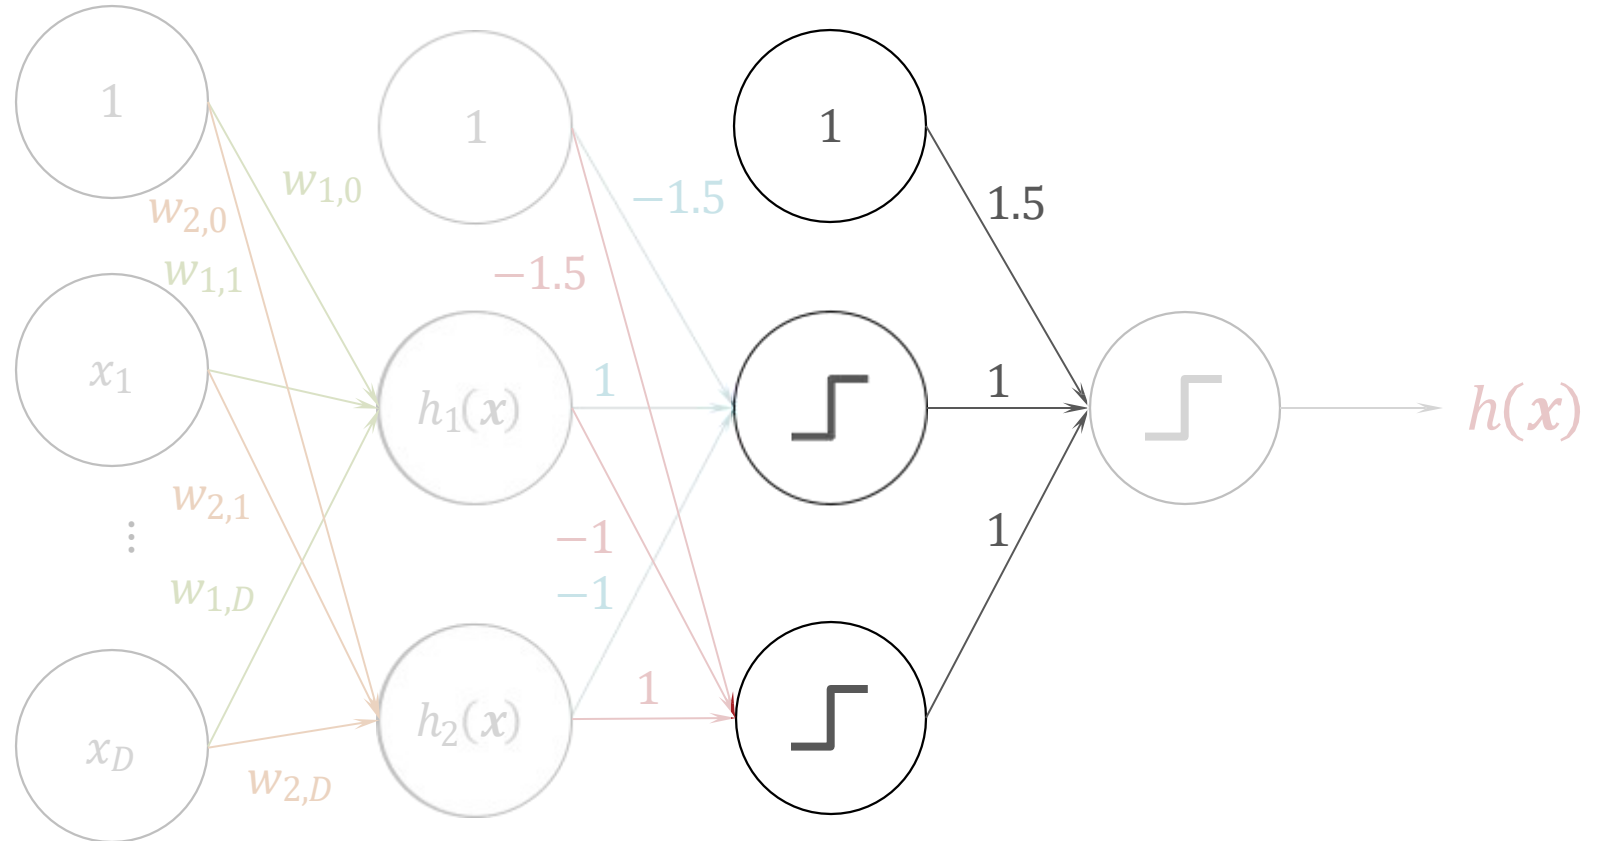
$$h(\mathbf{x}) = OR(AND(h_1(\mathbf{x}), \neg h_2(\mathbf{x})), AND(\neg h_1(\mathbf{x}), h_2(\mathbf{x})))$$



$$h(\mathbf{x}) = \text{sign}(\text{sign}(\text{sign}(\mathbf{w}_1^T \mathbf{x}) - \text{sign}(\mathbf{w}_2^T \mathbf{x}) - 1.5) + \text{sign}(-\text{sign}(\mathbf{w}_1^T \mathbf{x}) + \text{sign}(\mathbf{w}_2^T \mathbf{x}) - 1.5) + 1.5)$$

# Building a Network

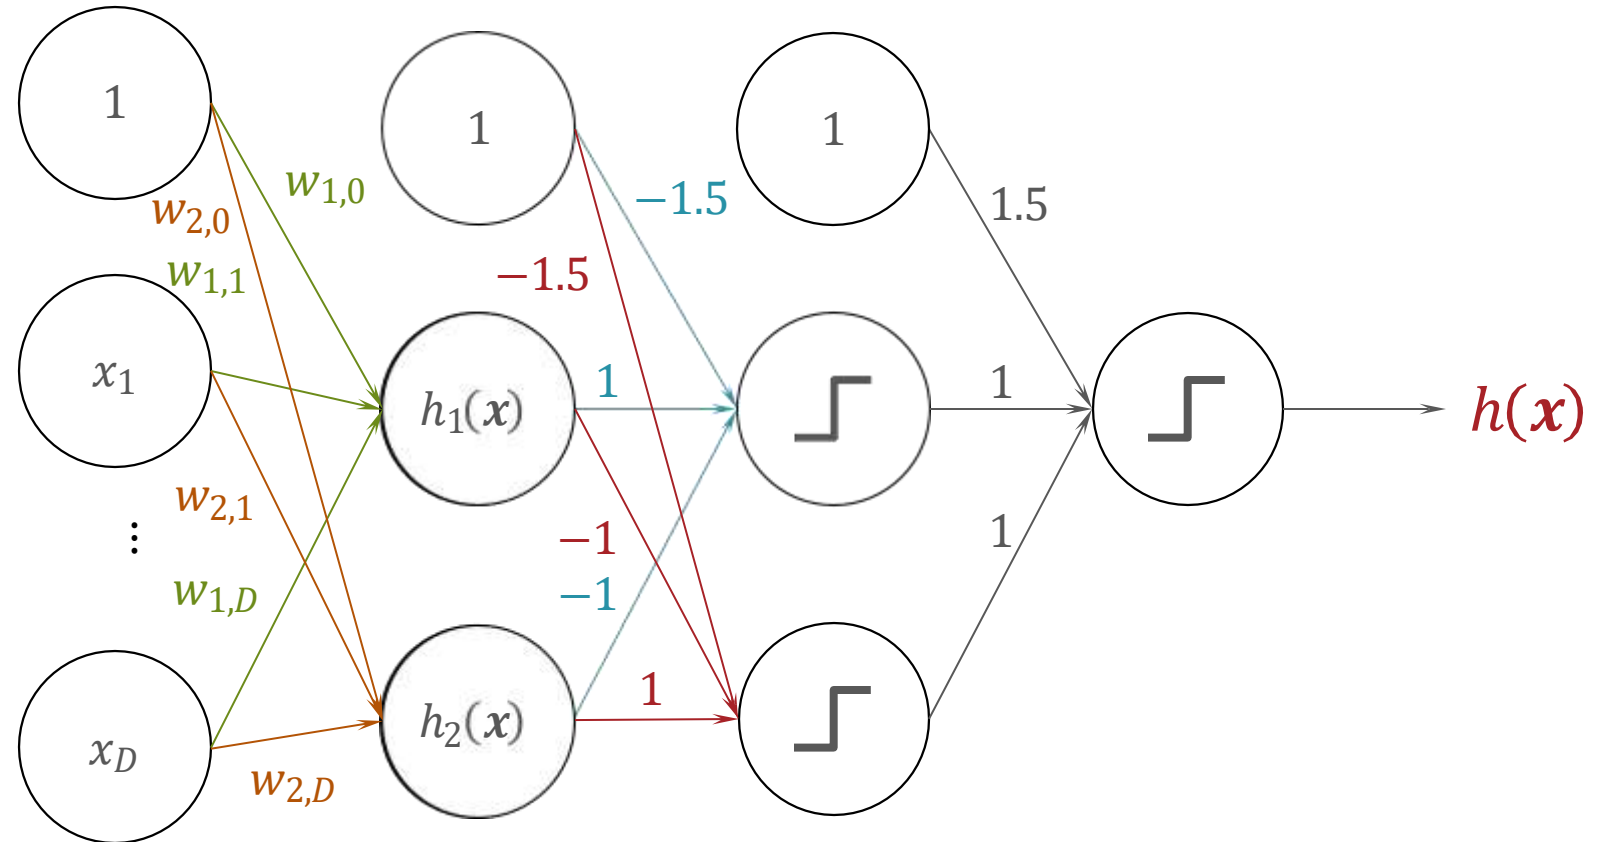
$$h(\mathbf{x}) = OR(AND(h_1(\mathbf{x}), \neg h_2(\mathbf{x})), AND(\neg h_1(\mathbf{x}), h_2(\mathbf{x})))$$



$$h(\mathbf{x}) = \text{sign}(\text{sign}(\text{sign}(\mathbf{w}_1^T \mathbf{x}) - \text{sign}(\mathbf{w}_2^T \mathbf{x}) - 1.5) + \text{sign}(-\text{sign}(\mathbf{w}_1^T \mathbf{x}) + \text{sign}(\mathbf{w}_2^T \mathbf{x}) - 1.5) + 1.5)$$

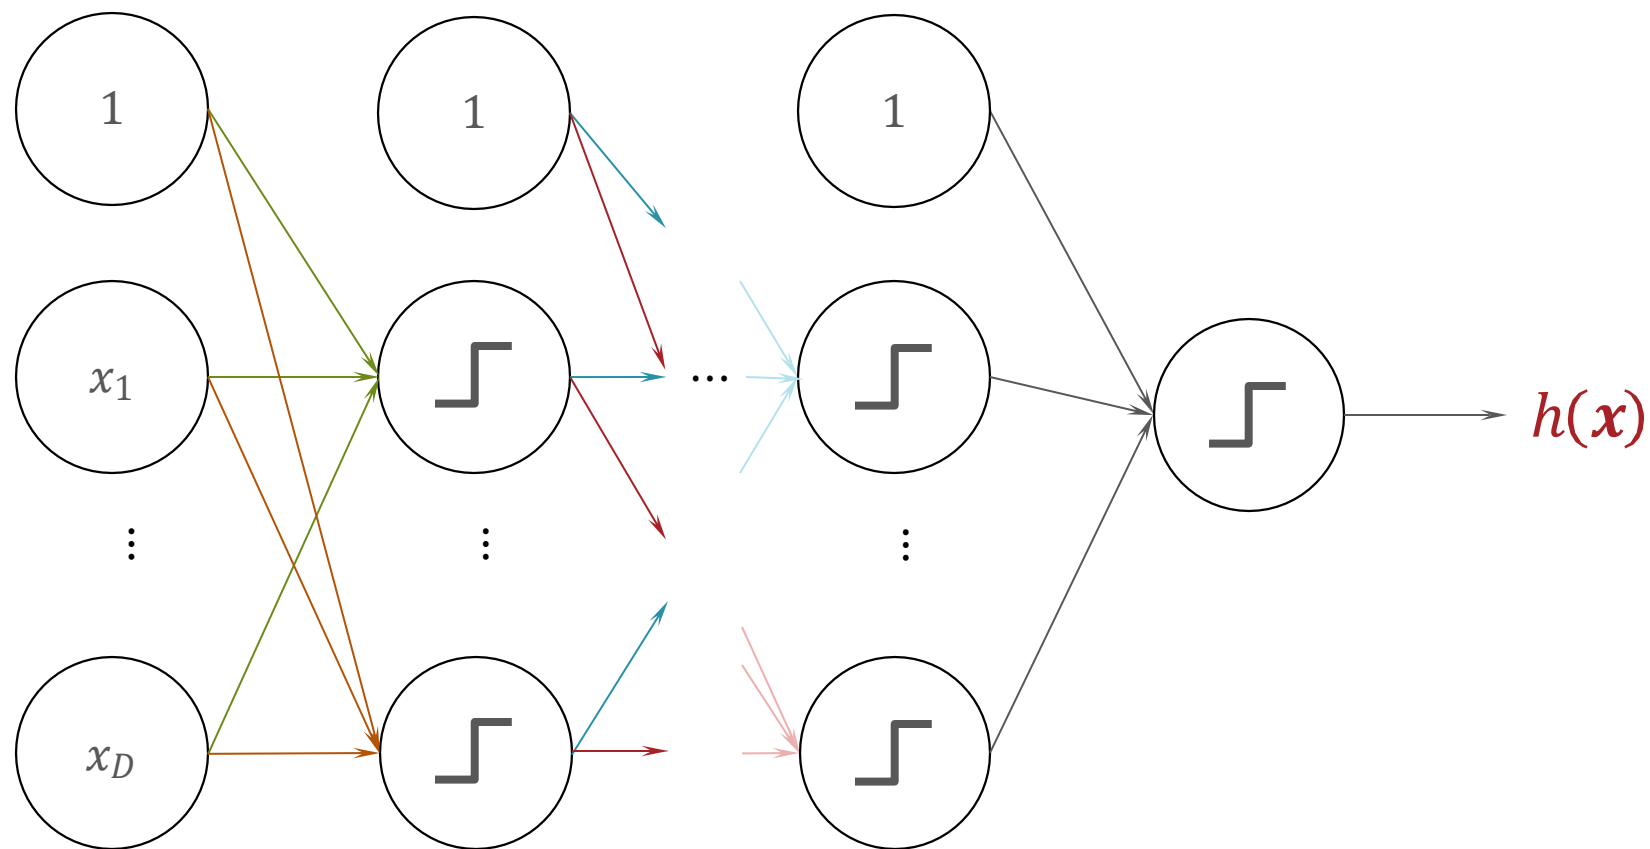
# Building a Network

$$h(\mathbf{x}) = OR(AND(h_1(\mathbf{x}), \neg h_2(\mathbf{x})), AND(\neg h_1(\mathbf{x}), h_2(\mathbf{x})))$$

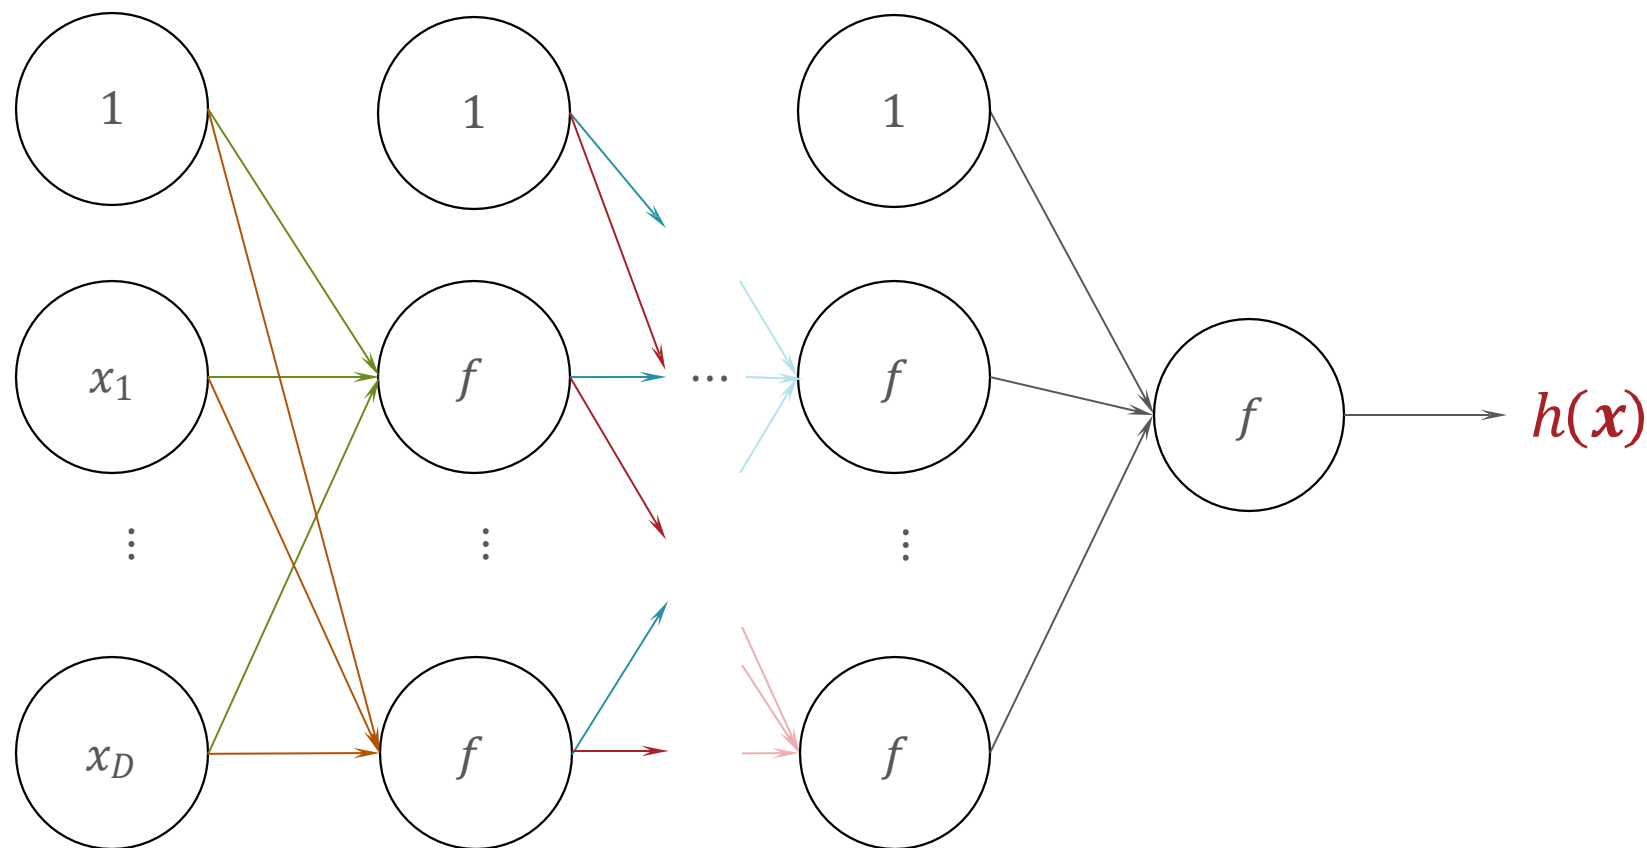


$$h(\mathbf{x}) = \text{sign}(\text{sign}(\text{sign}(\mathbf{w}_1^T \mathbf{x}) - \text{sign}(\mathbf{w}_2^T \mathbf{x}) - 1.5) + \text{sign}(-\text{sign}(\mathbf{w}_1^T \mathbf{x}) + \text{sign}(\mathbf{w}_2^T \mathbf{x}) - 1.5) + 1.5)$$

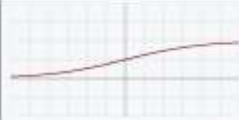
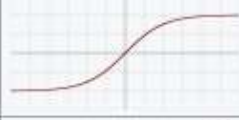






# Multi-Layer Perceptron (MLP)



# (Fully-Connected) Feed Forward Neural Network



# Activation Functions

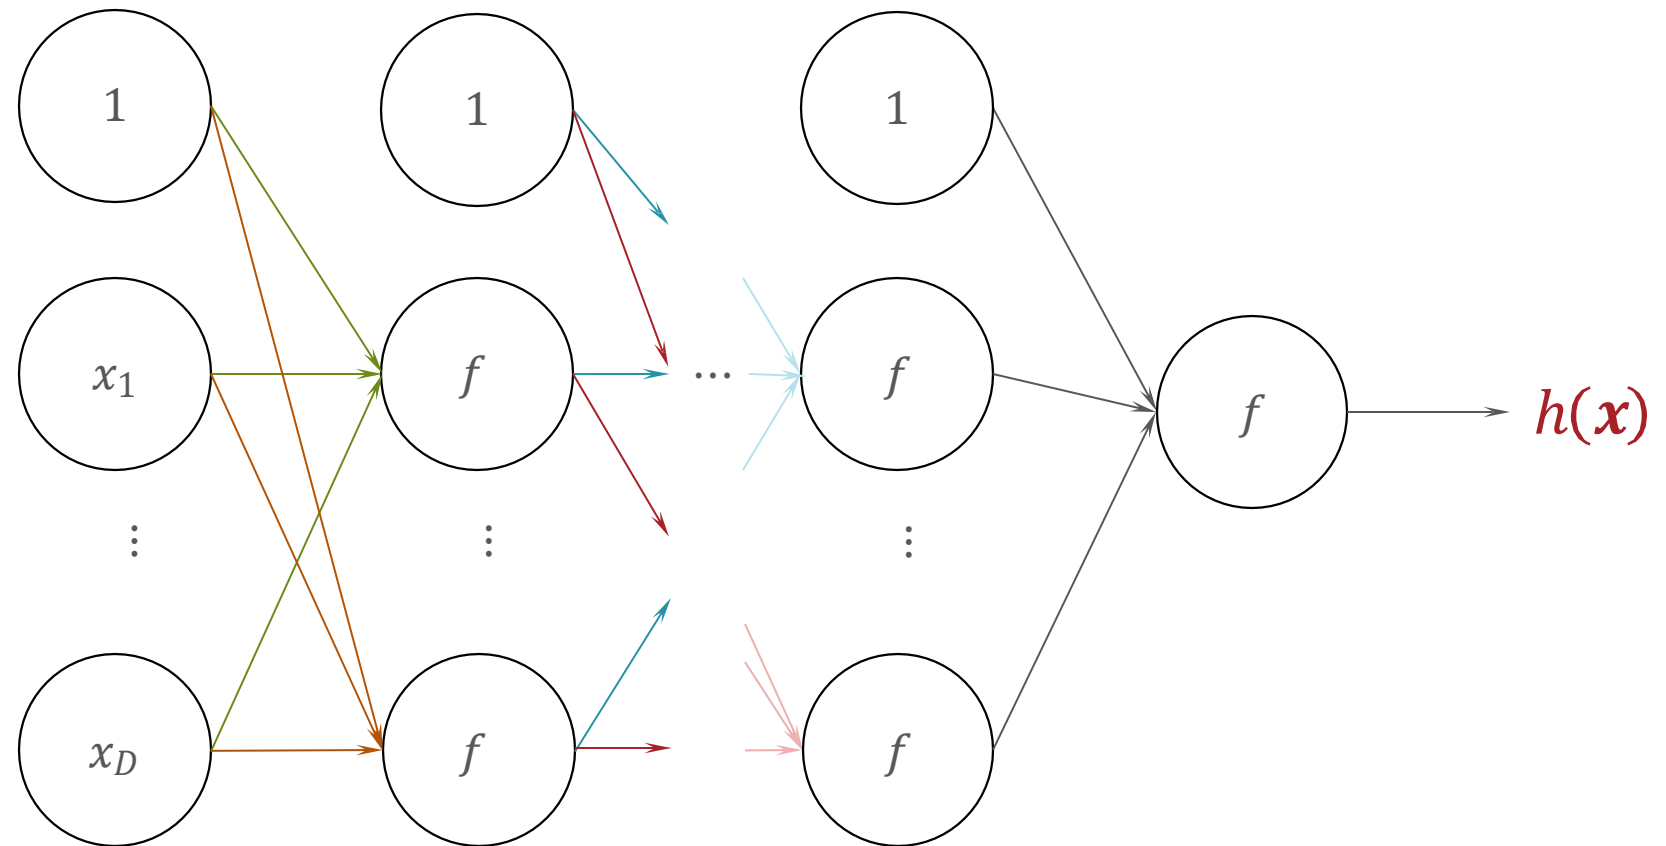
Logistic, sigmoid, or soft step		$\sigma(x) = \frac{1}{1 + e^{-x}}$
Hyperbolic tangent (tanh)		$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$
Rectified linear unit (ReLU) <sup>[7]</sup>		$\begin{cases} 0 & \text{if } x \leq 0 \\ x & \text{if } x > 0 \end{cases}$ $= \max\{0, x\} = x \mathbf{1}_{x>0}$
Gaussian Error Linear Unit (GELU) <sup>[4]</sup>		$\frac{1}{2}x \left( 1 + \operatorname{erf}\left(\frac{x}{\sqrt{2}}\right) \right)$ $= x\Phi(x)$
Softplus <sup>[8]</sup>		$\ln(1 + e^x)$
Exponential linear unit (ELU) <sup>[9]</sup>		$\begin{cases} \alpha(e^x - 1) & \text{if } x \leq 0 \\ x & \text{if } x > 0 \end{cases}$ <p>with parameter <math>\alpha</math></p>
Leaky rectified linear unit (Leaky ReLU) <sup>[11]</sup>		$\begin{cases} 0.01x & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$
Parametric rectified linear unit (PReLU) <sup>[12]</sup>		$\begin{cases} \alpha x & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$ <p>with parameter <math>\alpha</math></p>



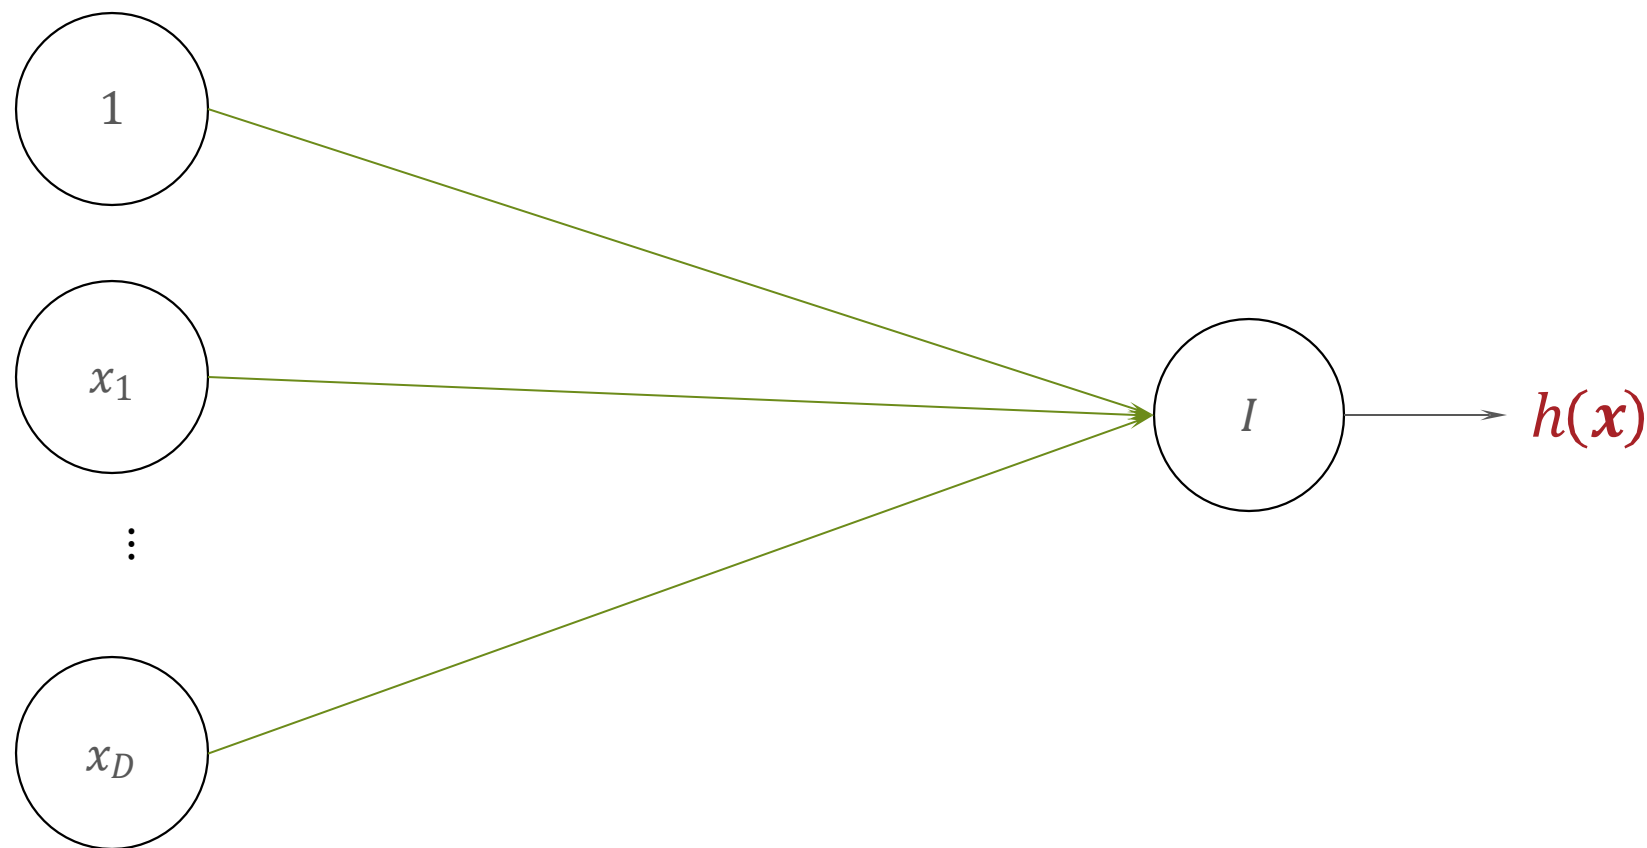
## Poll Question 1

True or False: Linear and logistic regression models can be expressed as neural networks.

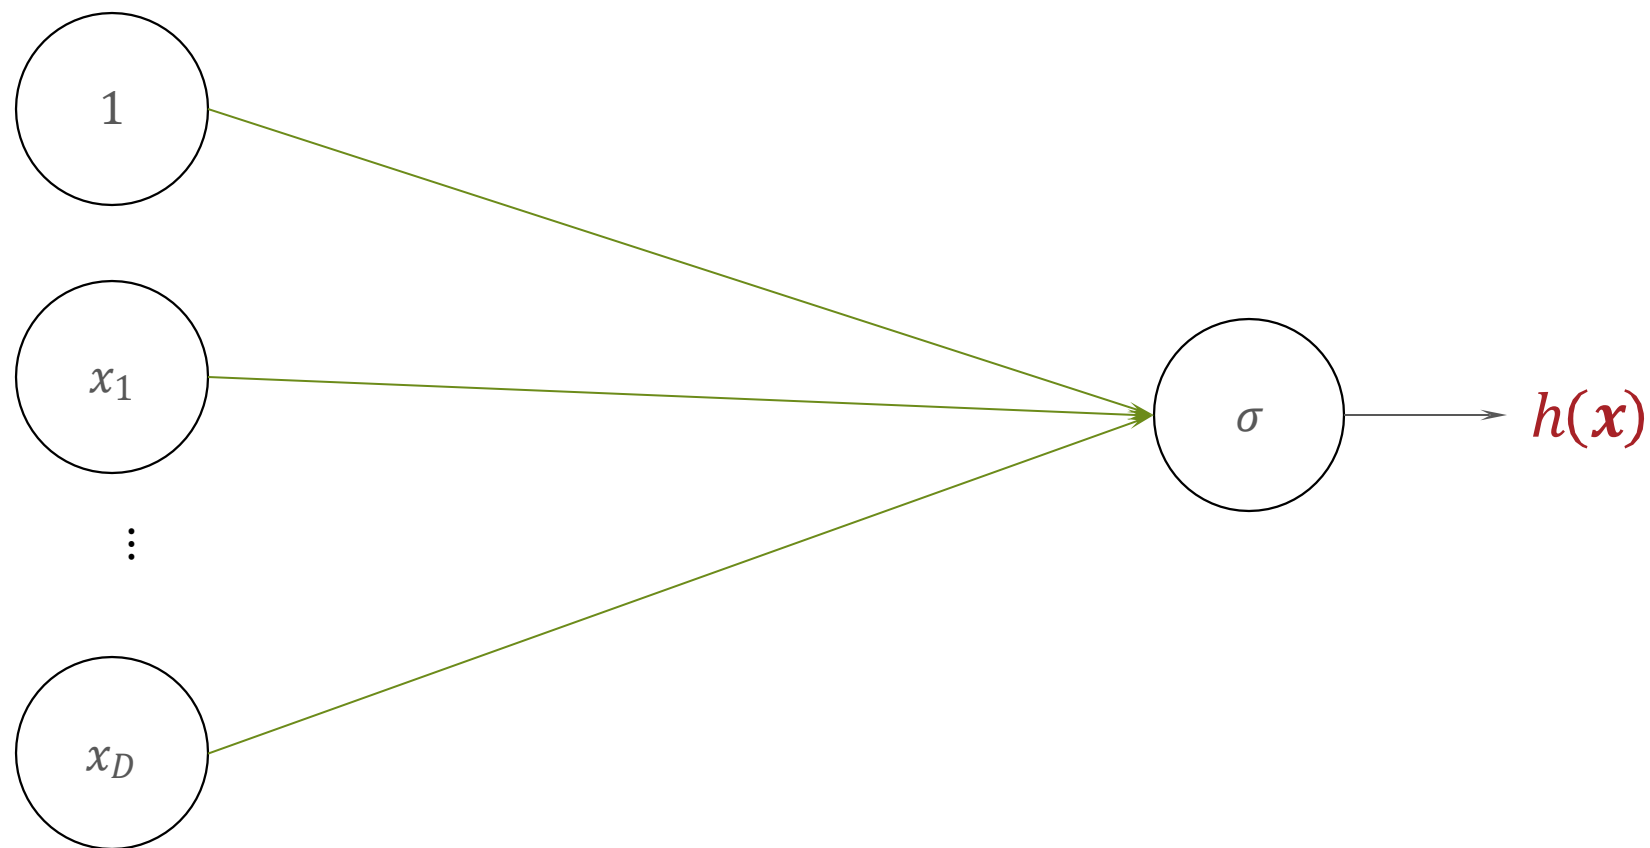
- A. Only true for linear regression
- B. Only true for logistic regression
- C. TOXIC
- D. True for both
- E. False for both



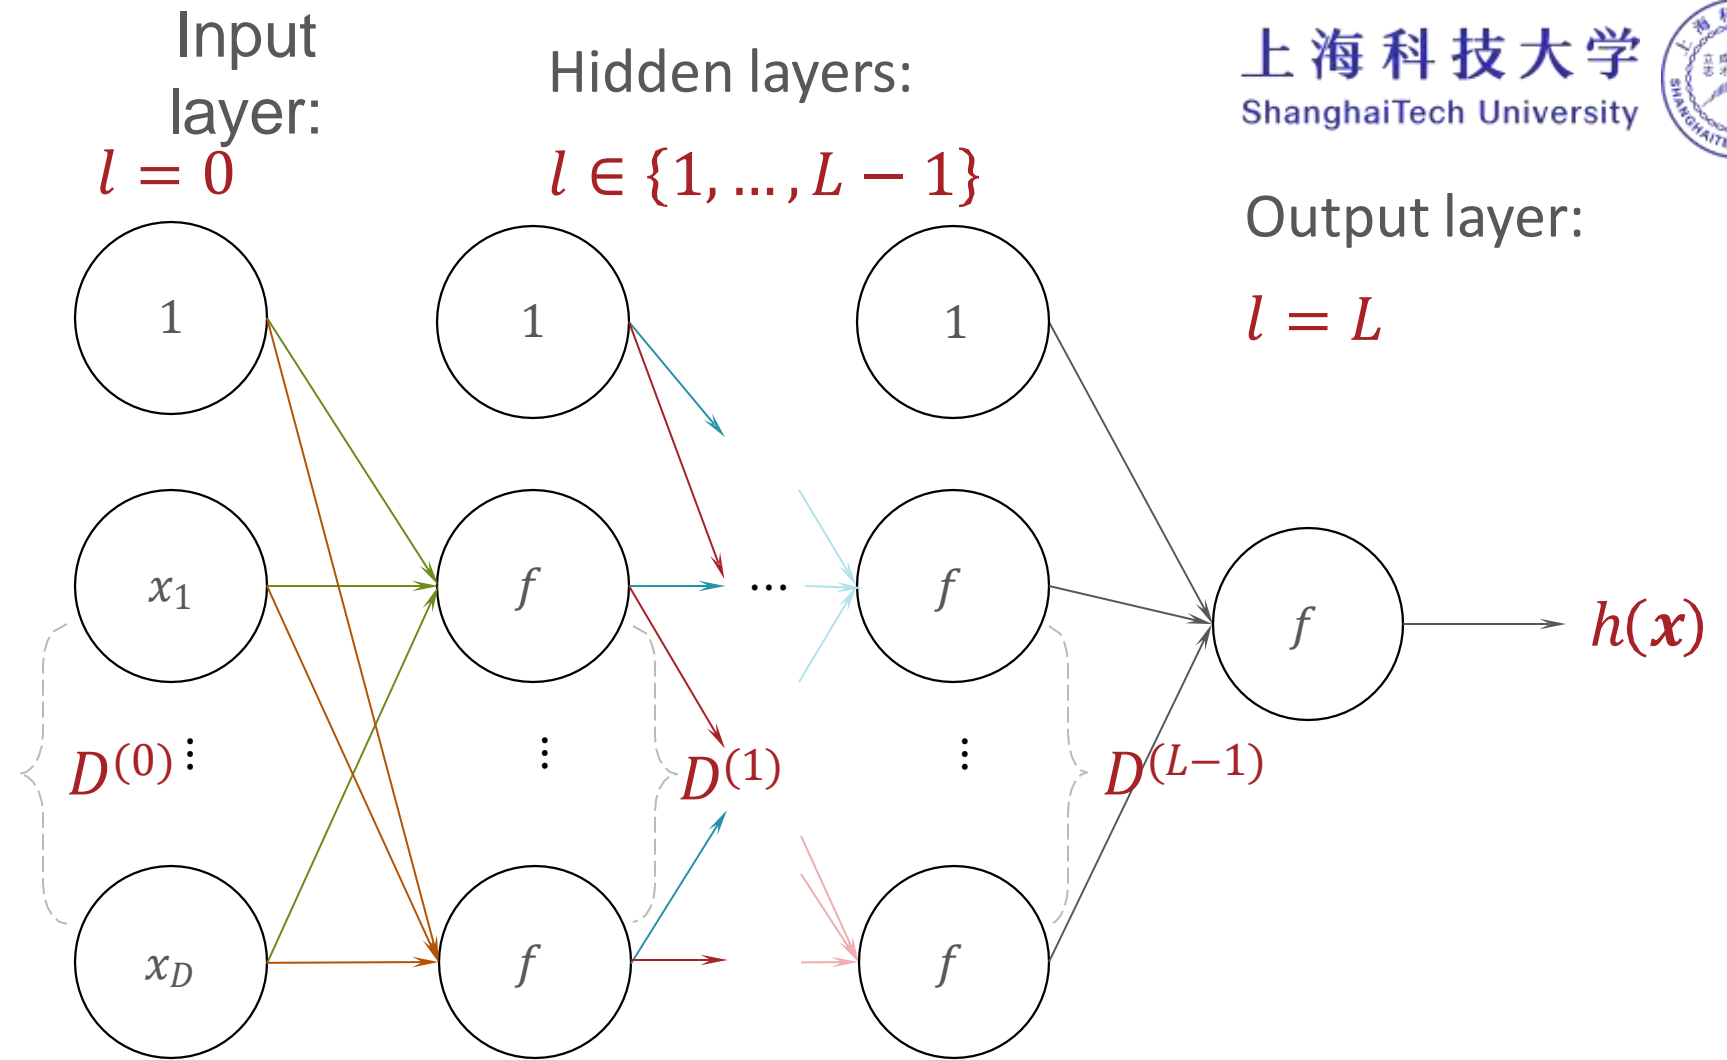
# Linear Regression as a Neural Network



# Logistic Regression as a Neural Network



# (Fully-Connected) Feed Forward Neural Network



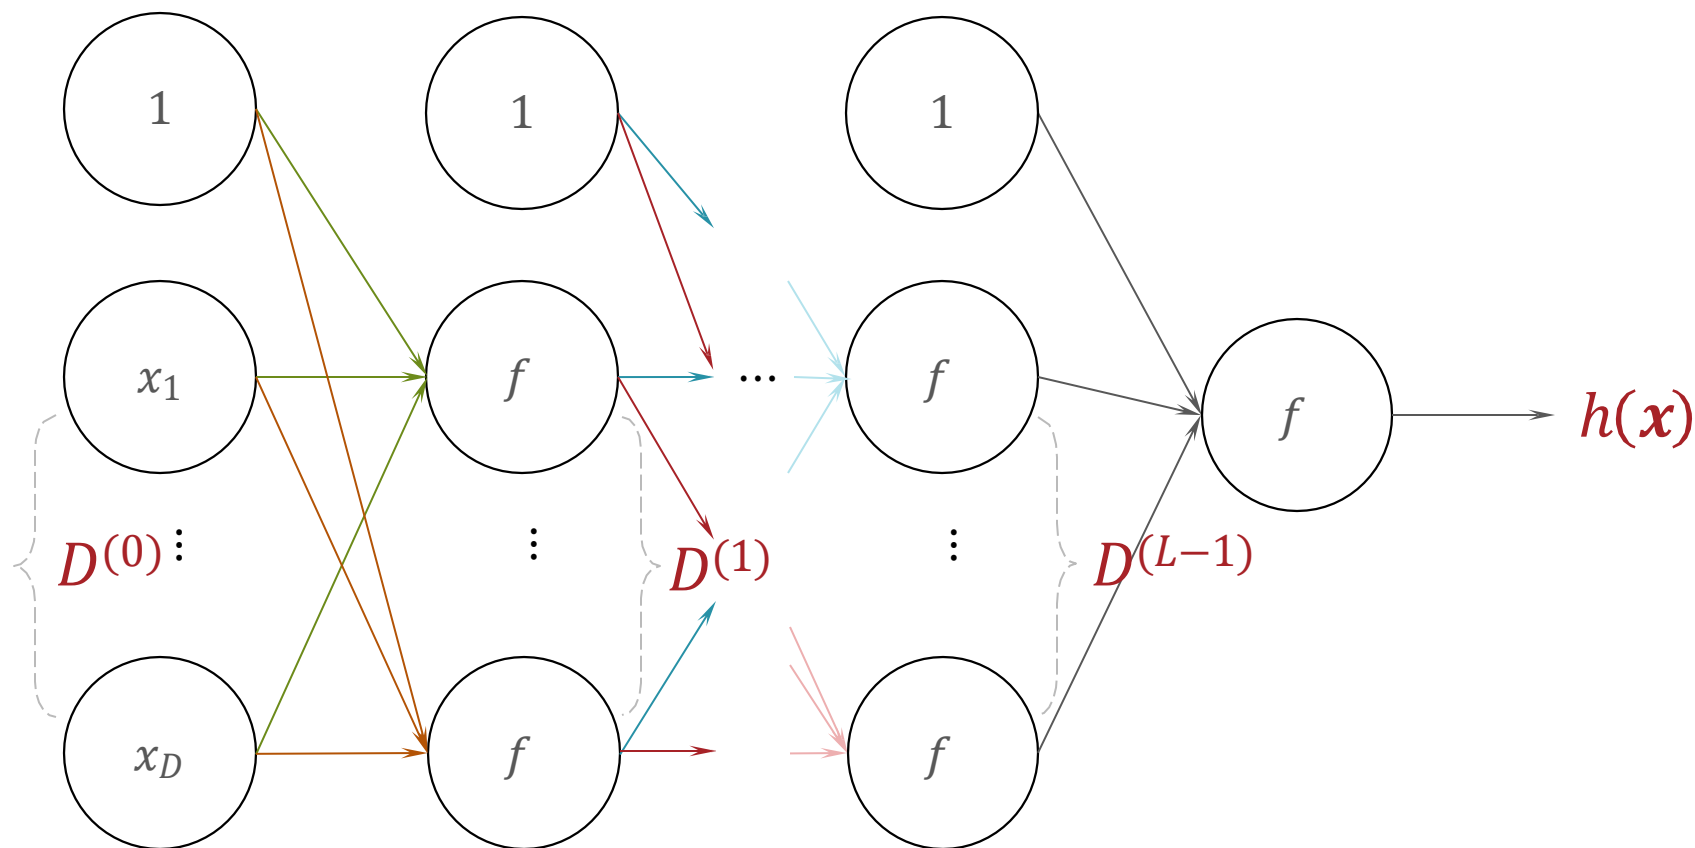
Layer  $l$  has dimension  $D^{(l)}$  → Layer  $l$  has  $D^{(l)} + 1$  nodes, counting the bias node

The weights between layer  $l - 1$  and layer  $l$  are a matrix:



# (Fully-Connected) Feed Forward Neural Network

$$W^{(l)} \in \mathbb{R}^{D^{(l)} \times (D^{(l-1)}+1)}$$

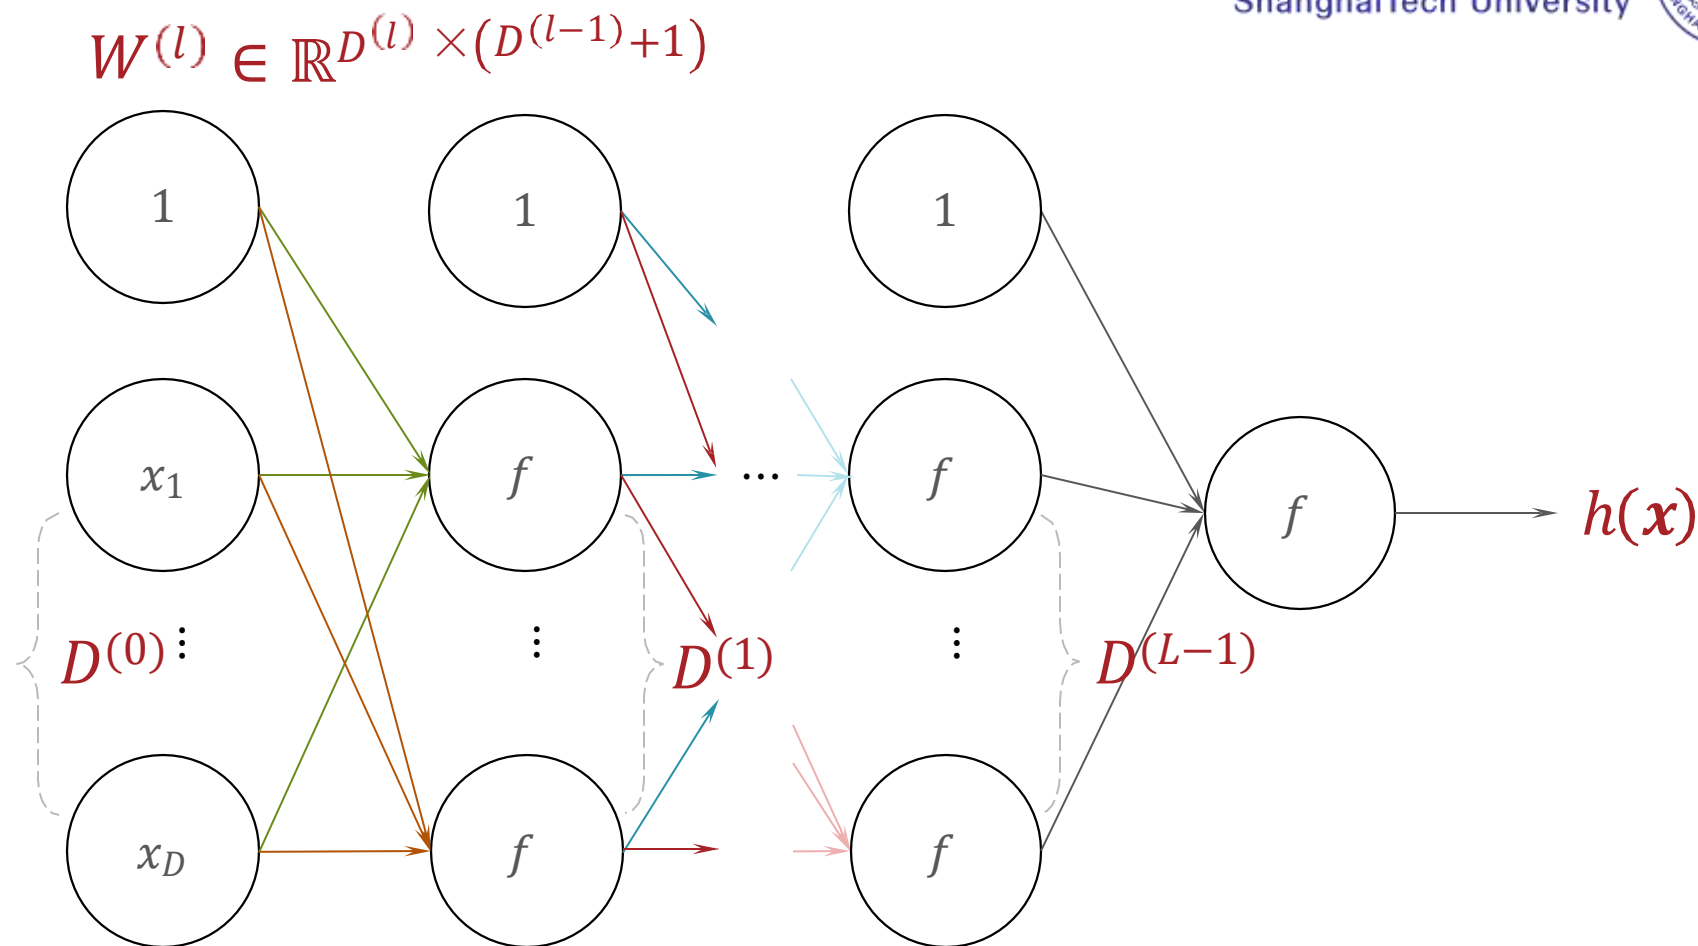


$w_{j,i}^{(l)}$  is the weight between node  $i$  in layer  $l - 1$  and node  $j$  in layer  $l$

The weights between layer  $l - 1$  and layer  $l$  are a matrix:



So what are all these layers doing for us anyway?



$w_{j,i}^{(l)}$  is the weight between node  $i$  in layer  $l - 1$  and node  $j$  in layer  $l$

# Neural Network Decision Boundaries: Example 1

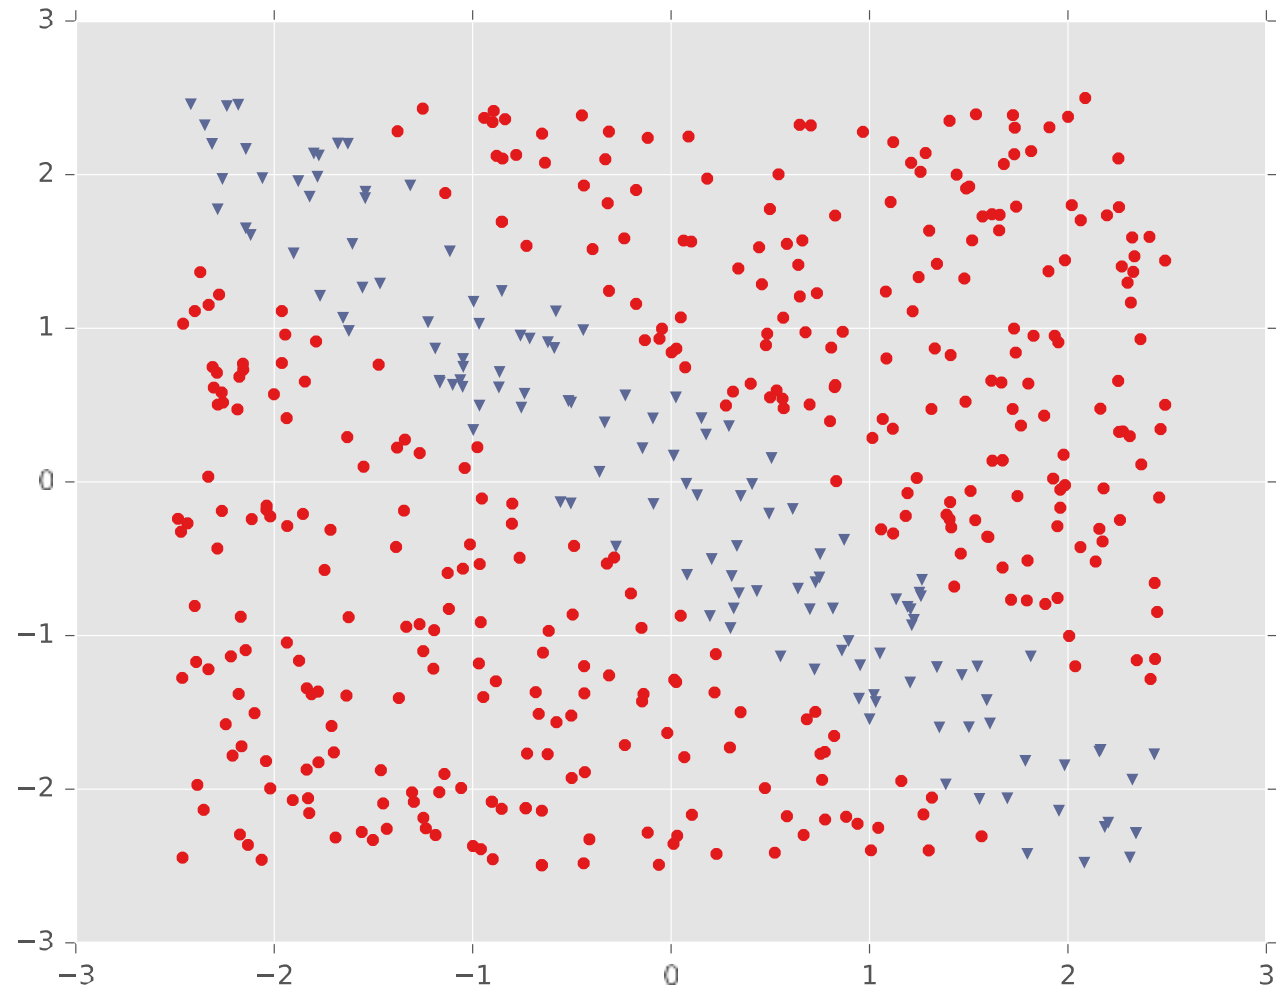


Figure courtesy of Matt Gormley

# Neural Network Decision Boundaries: Example 1

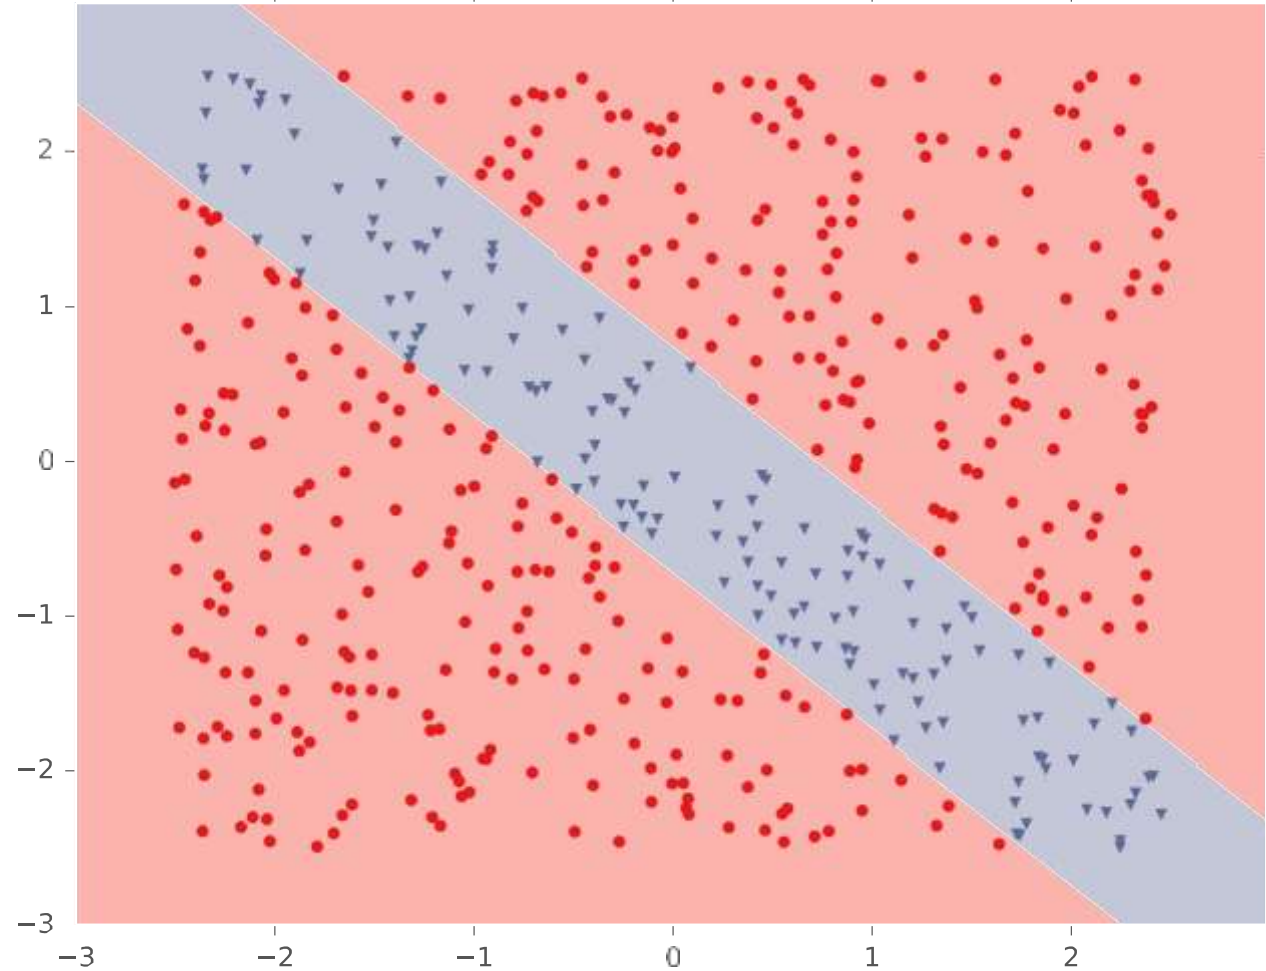


Figure courtesy of Matt Gormley

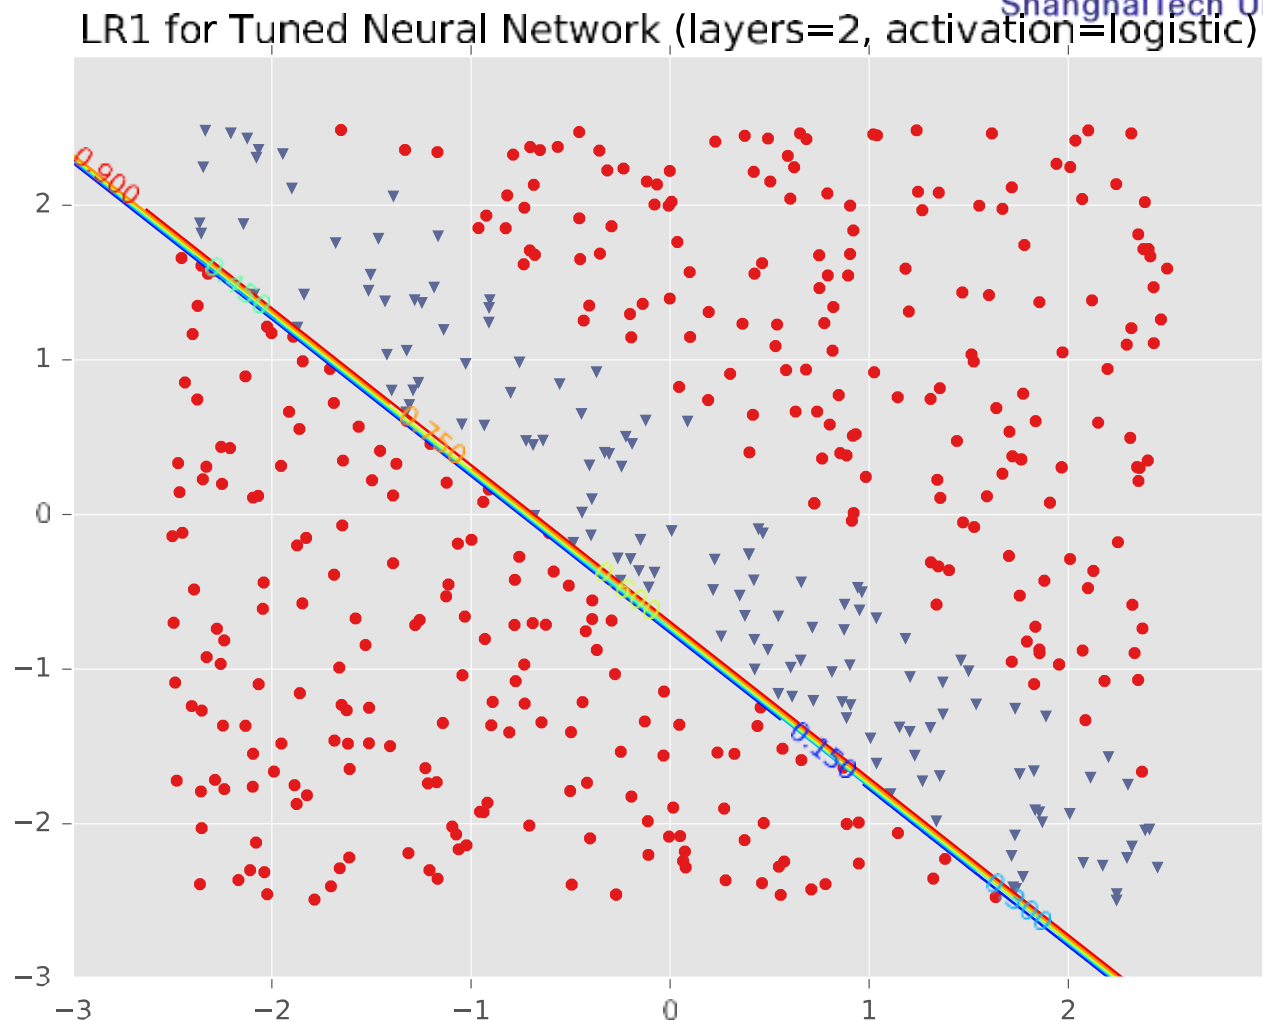


# Neural Network Decision Boundaries: Example 1

Tuned Neural Network (layers=2, activation=logistic)



# Neural Network Decision Boundaries: Example 1



# Neural Network Decision Boundaries: Example 1

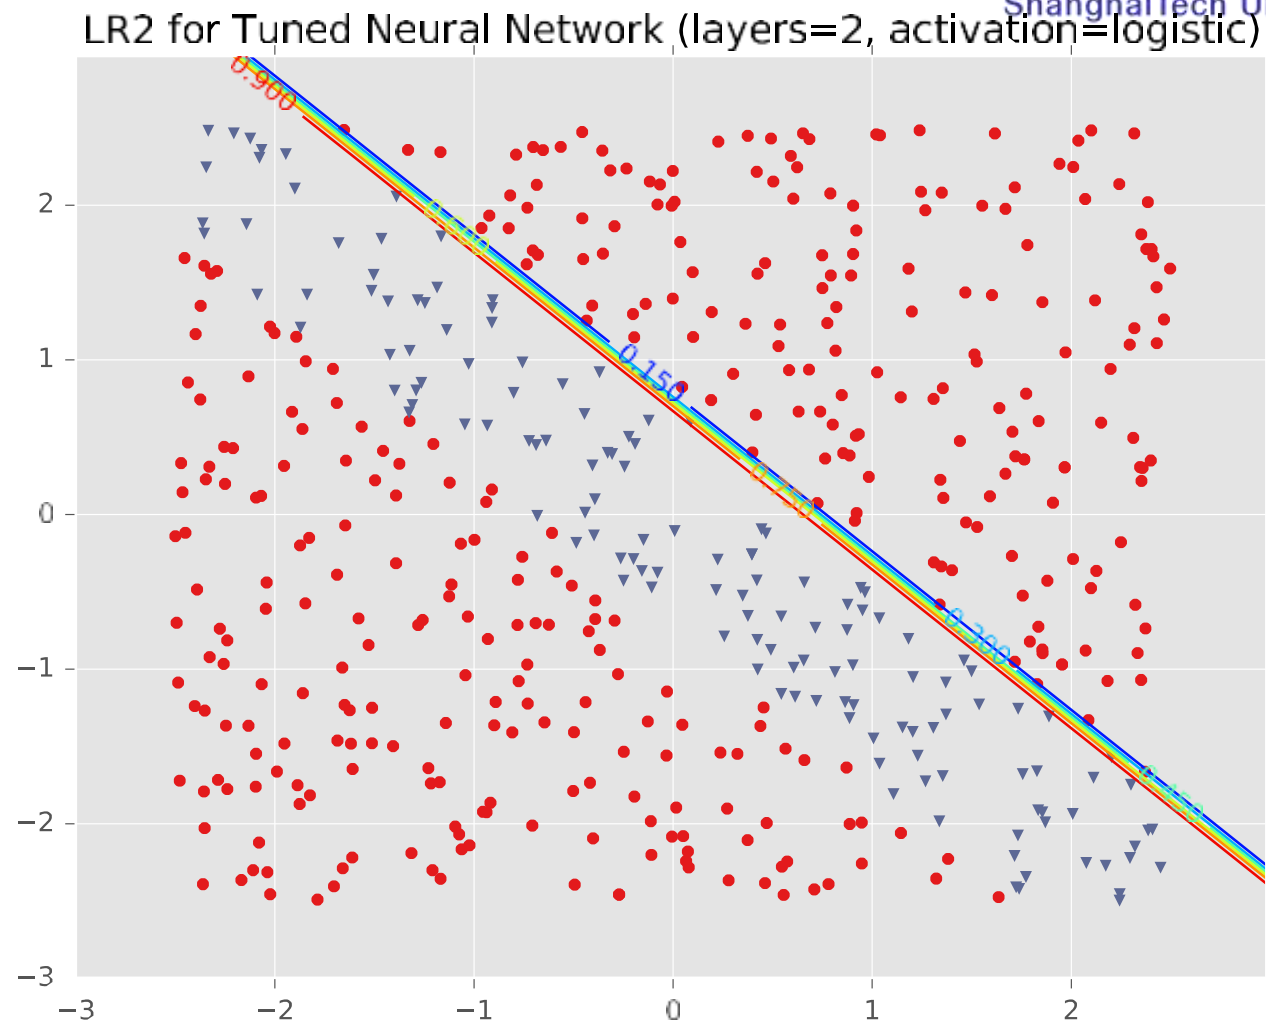


Figure courtesy of Matt Gormley

# Neural Network Decision Boundaries: Example 1

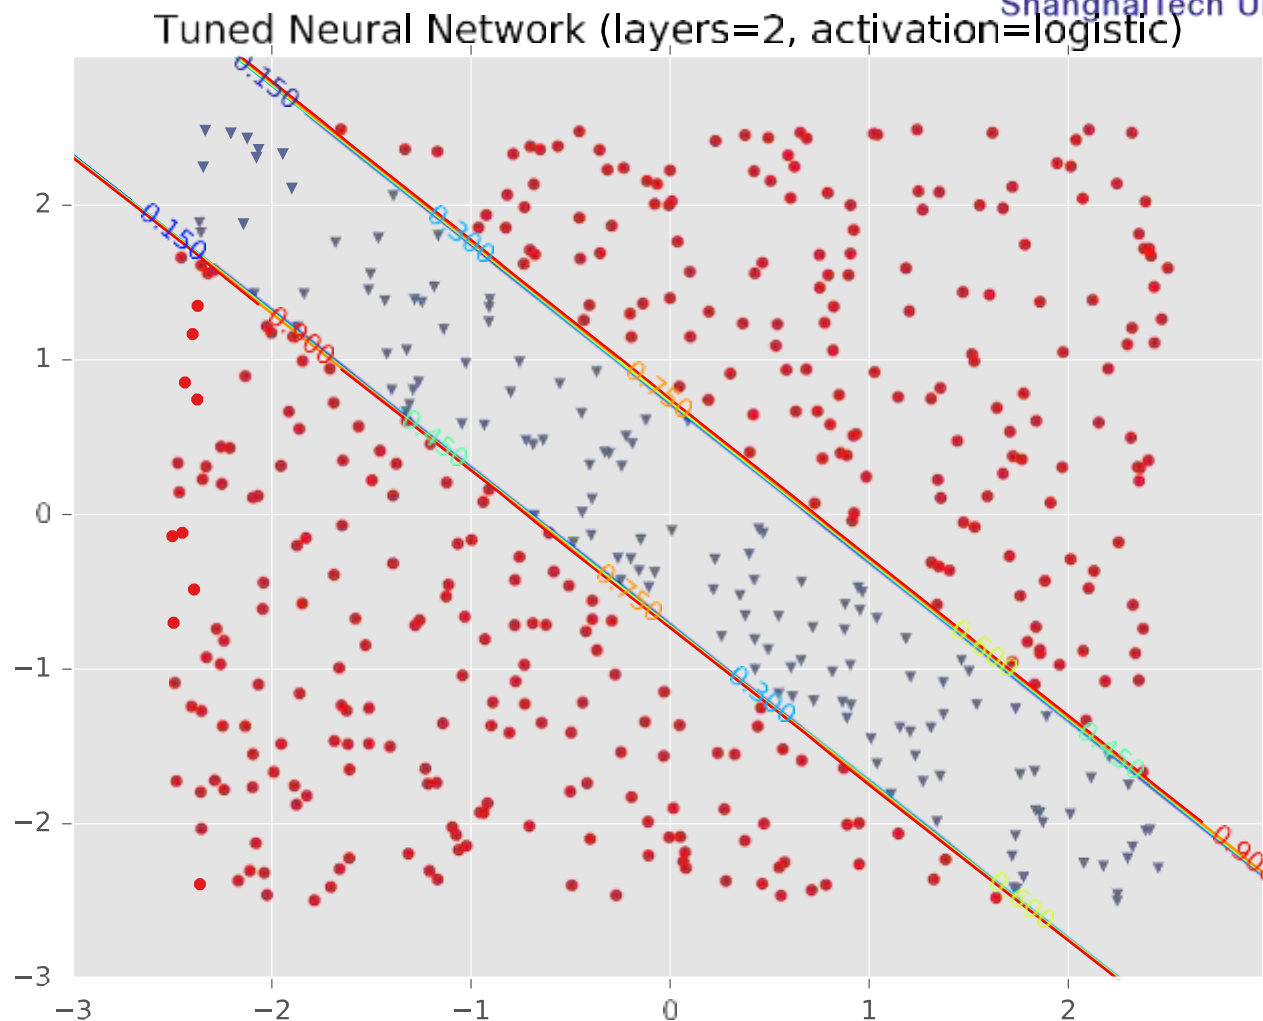
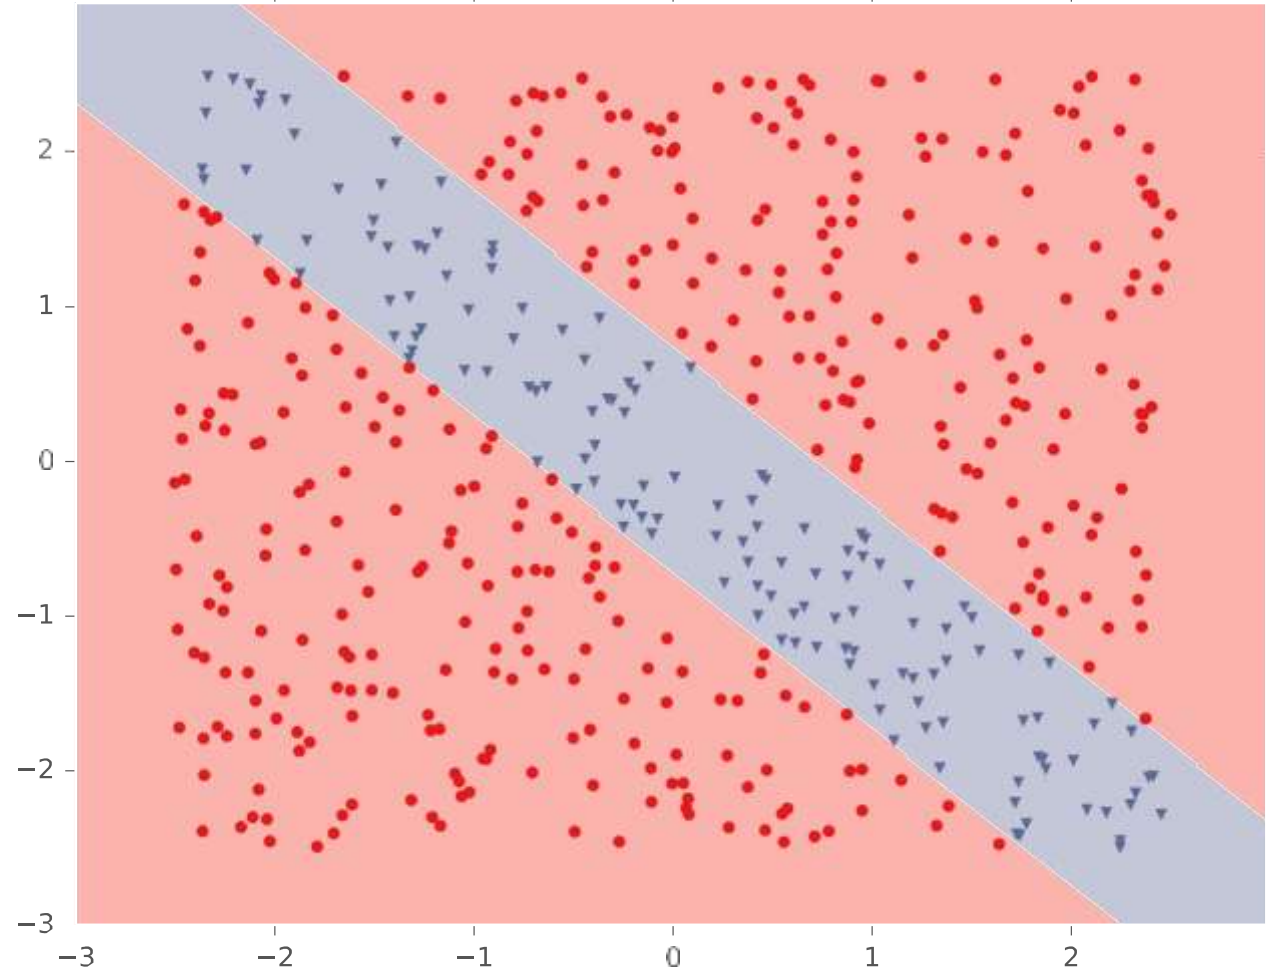


Figure courtesy of Matt Gormley

# Neural Network Decision Boundaries: Example 1

Tuned Neural Network (layers=2, activation=logistic)



# Neural Network Decision Boundaries: Example 2

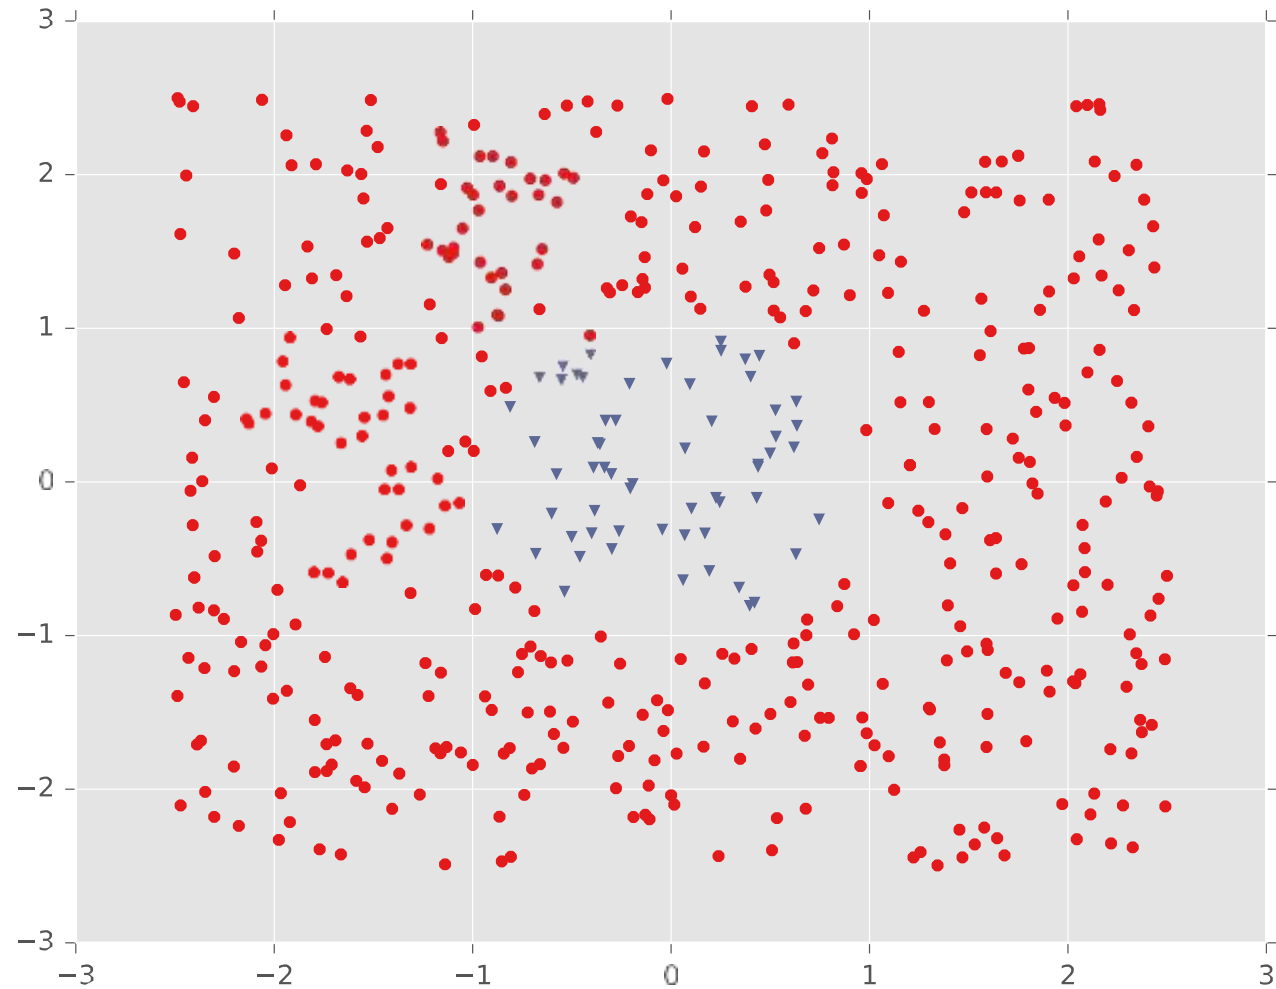


Figure courtesy of Matt Gormley

# Neural Network Decision Boundaries: Example 2

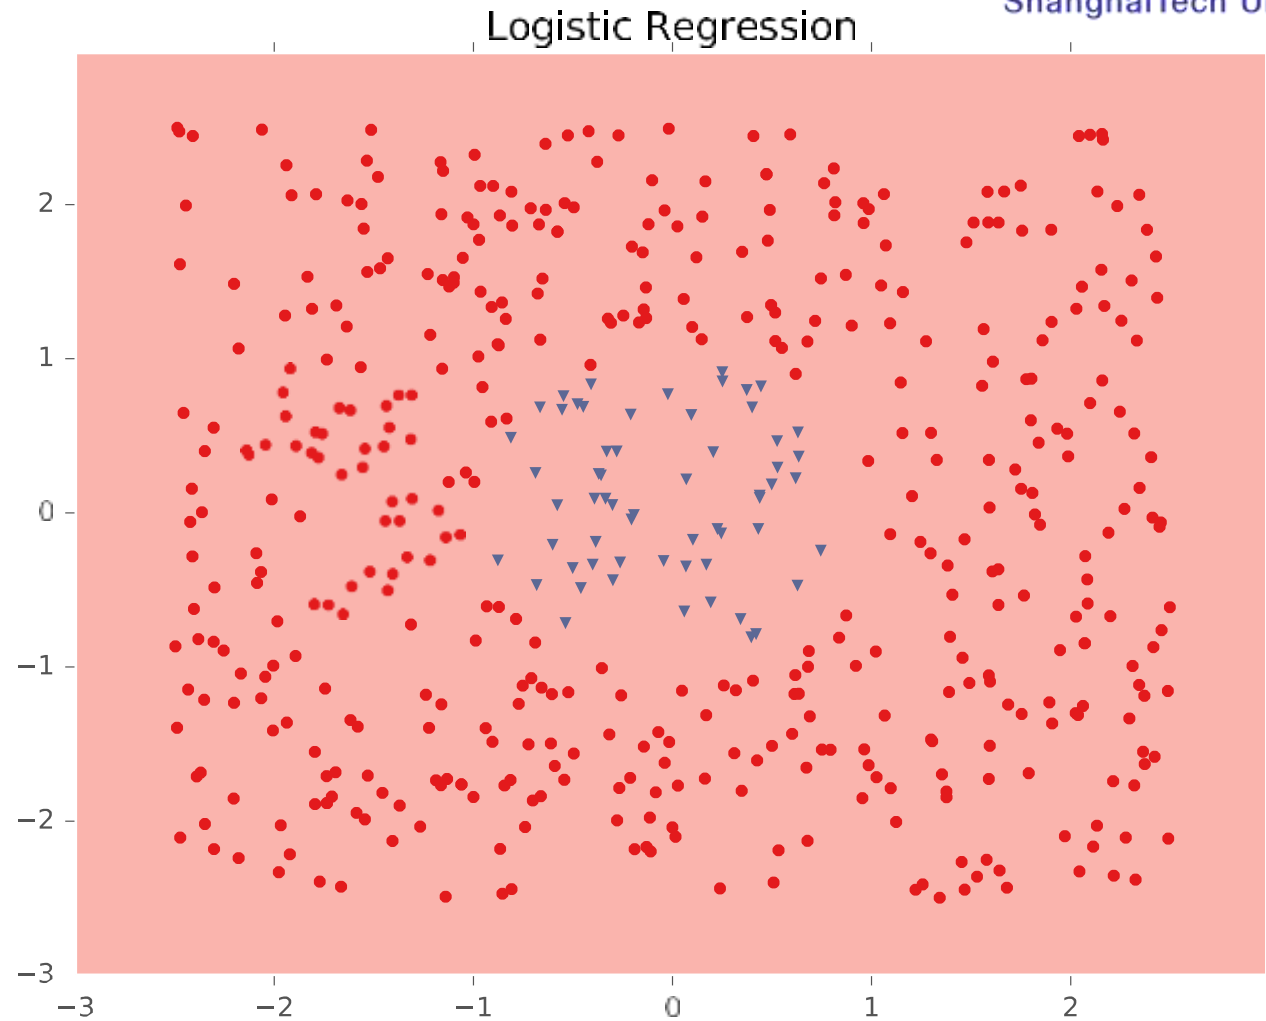
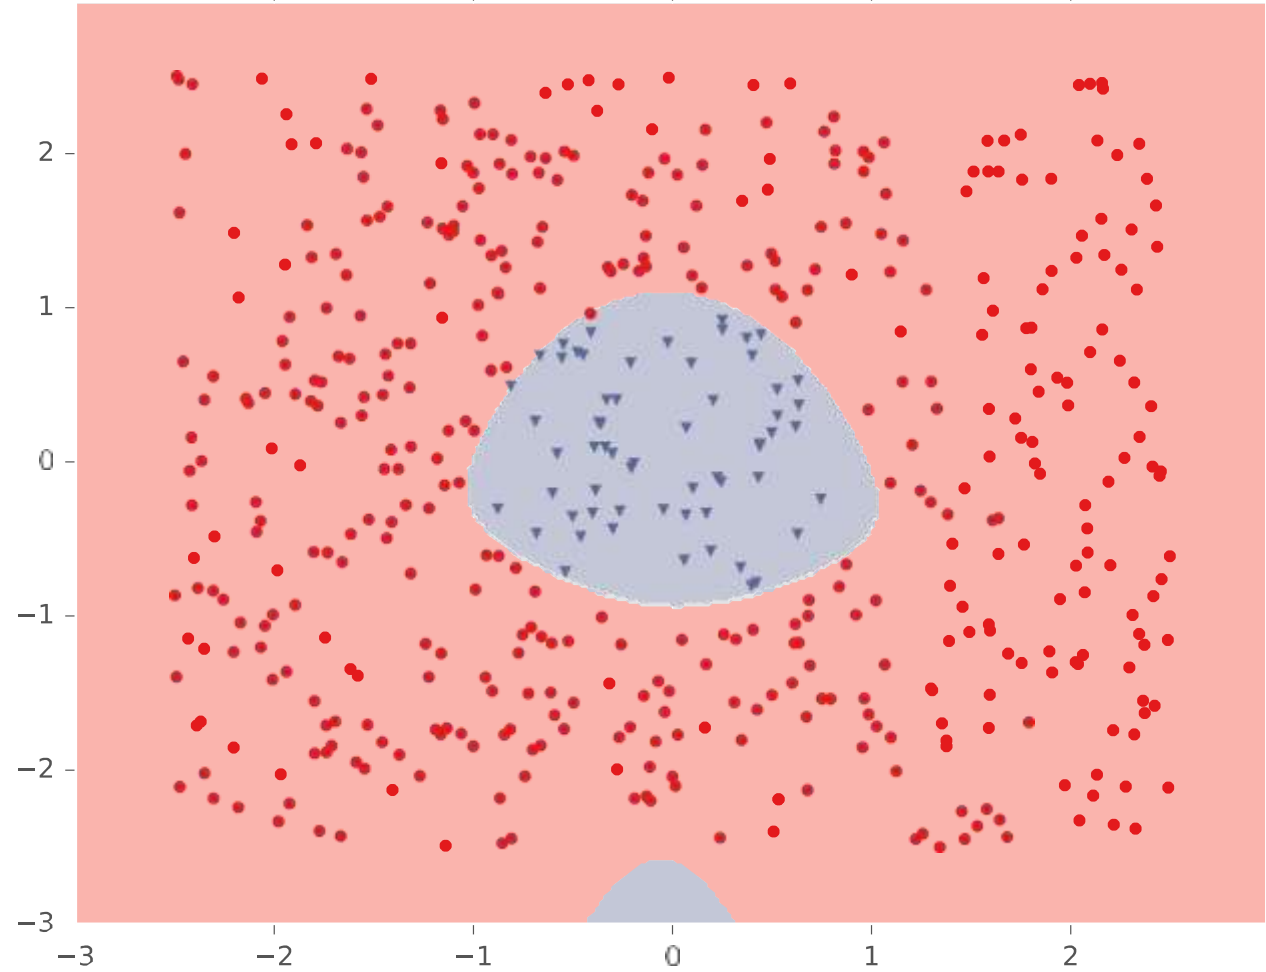


Figure courtesy of Matt Gormley

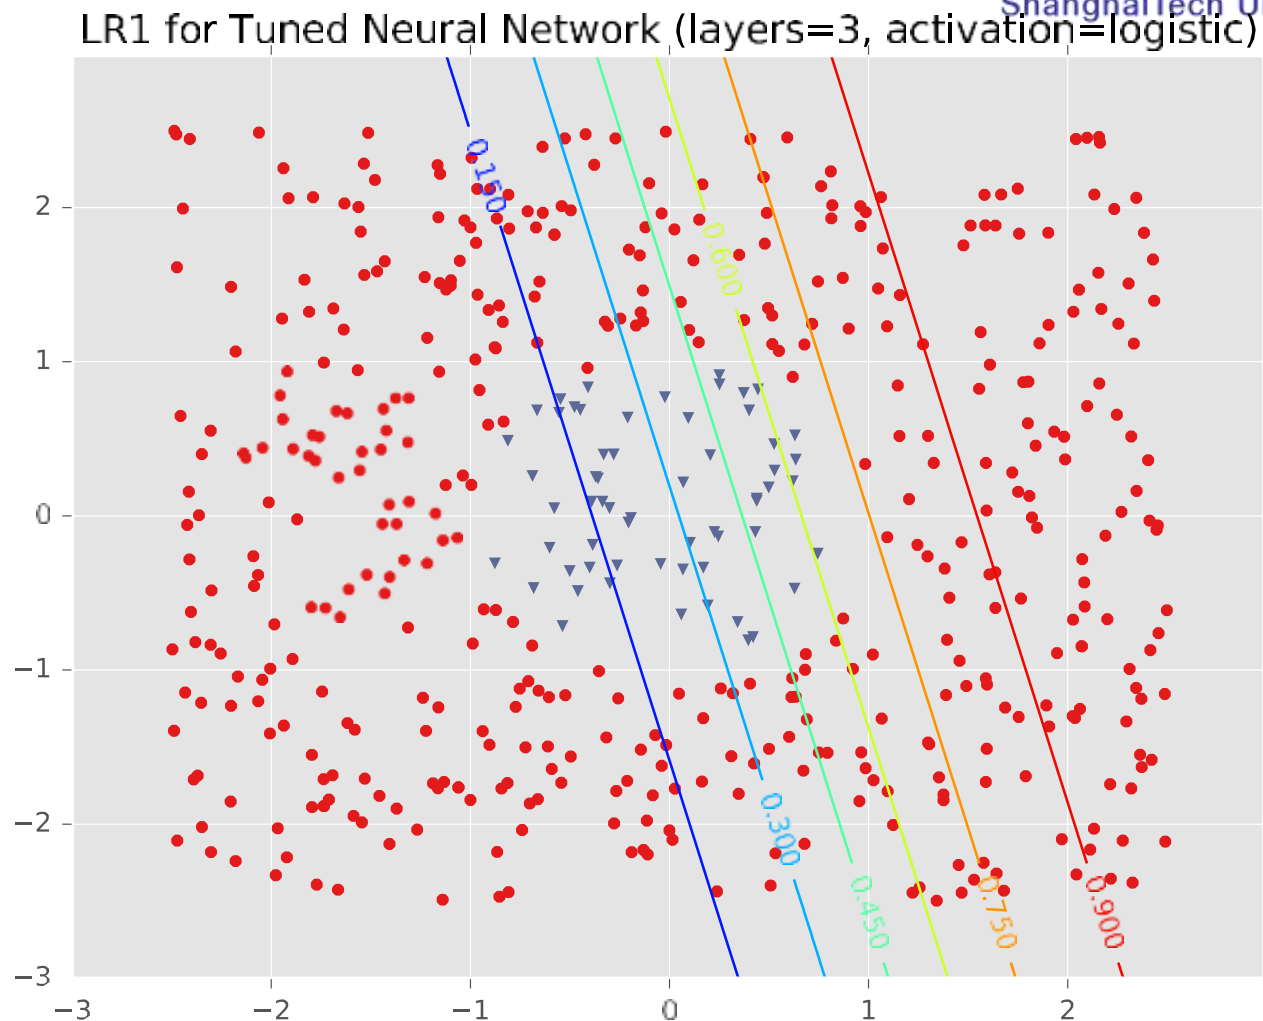
# Neural Network Decision Boundaries: Example 2

Tuned Neural Network (layers=3, activation=logistic)





# Neural Network Decision Boundaries: Example 2



# Neural Network Decision Boundaries: Example 2

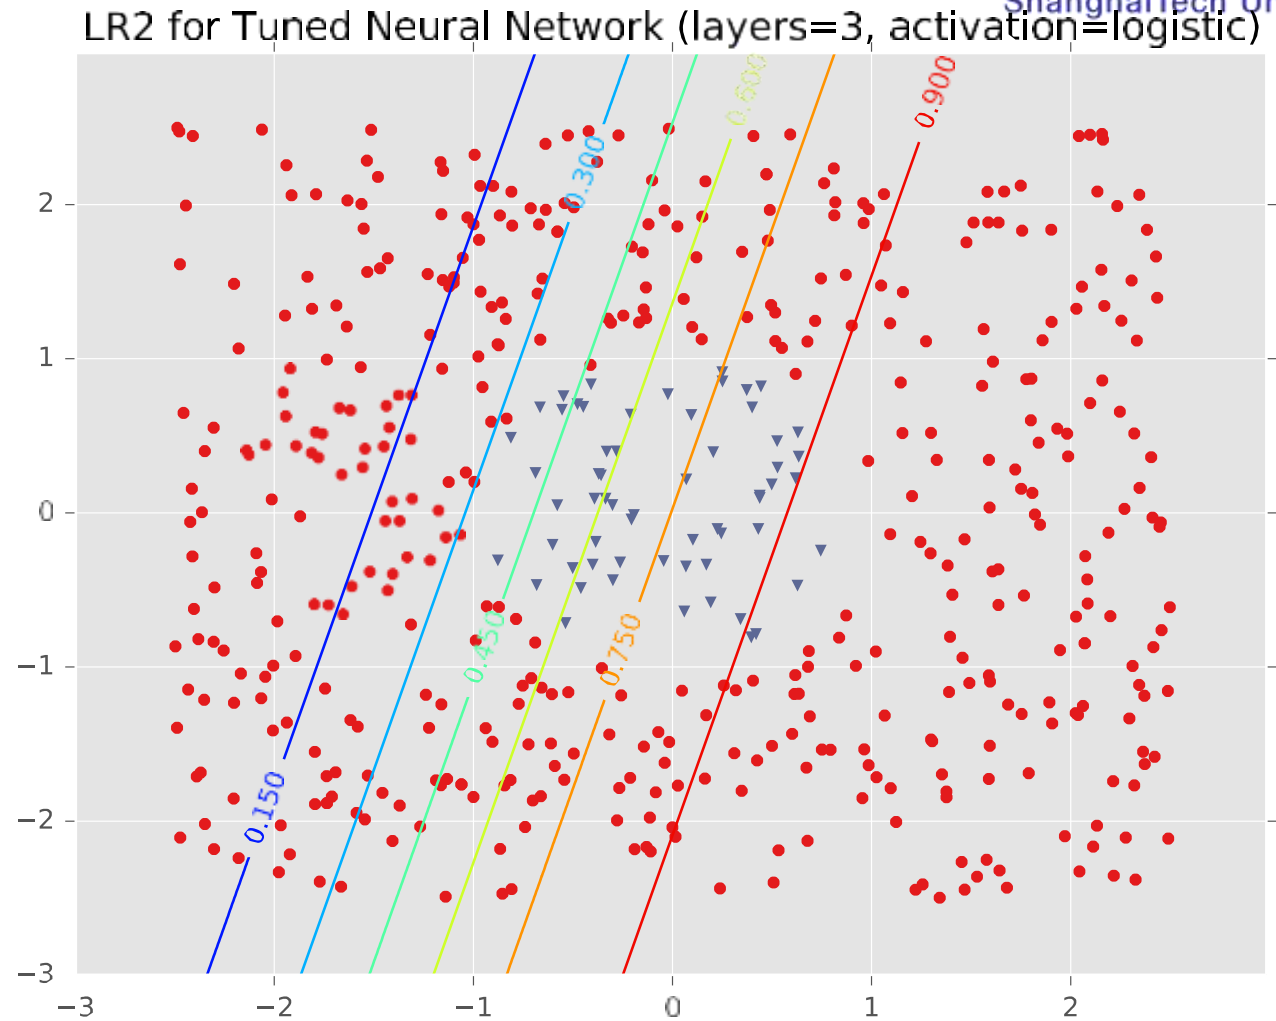
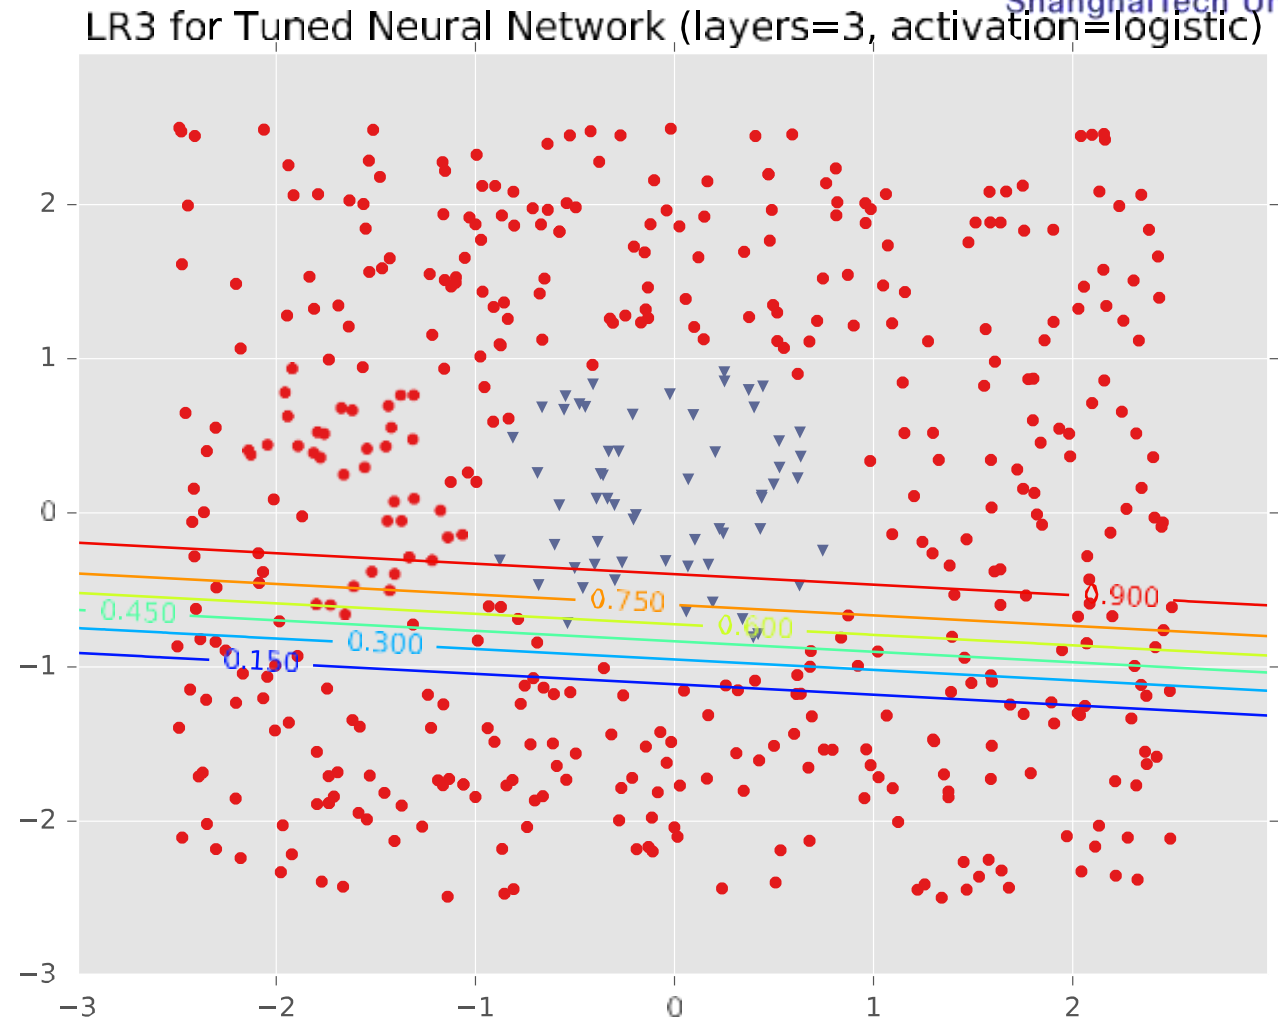
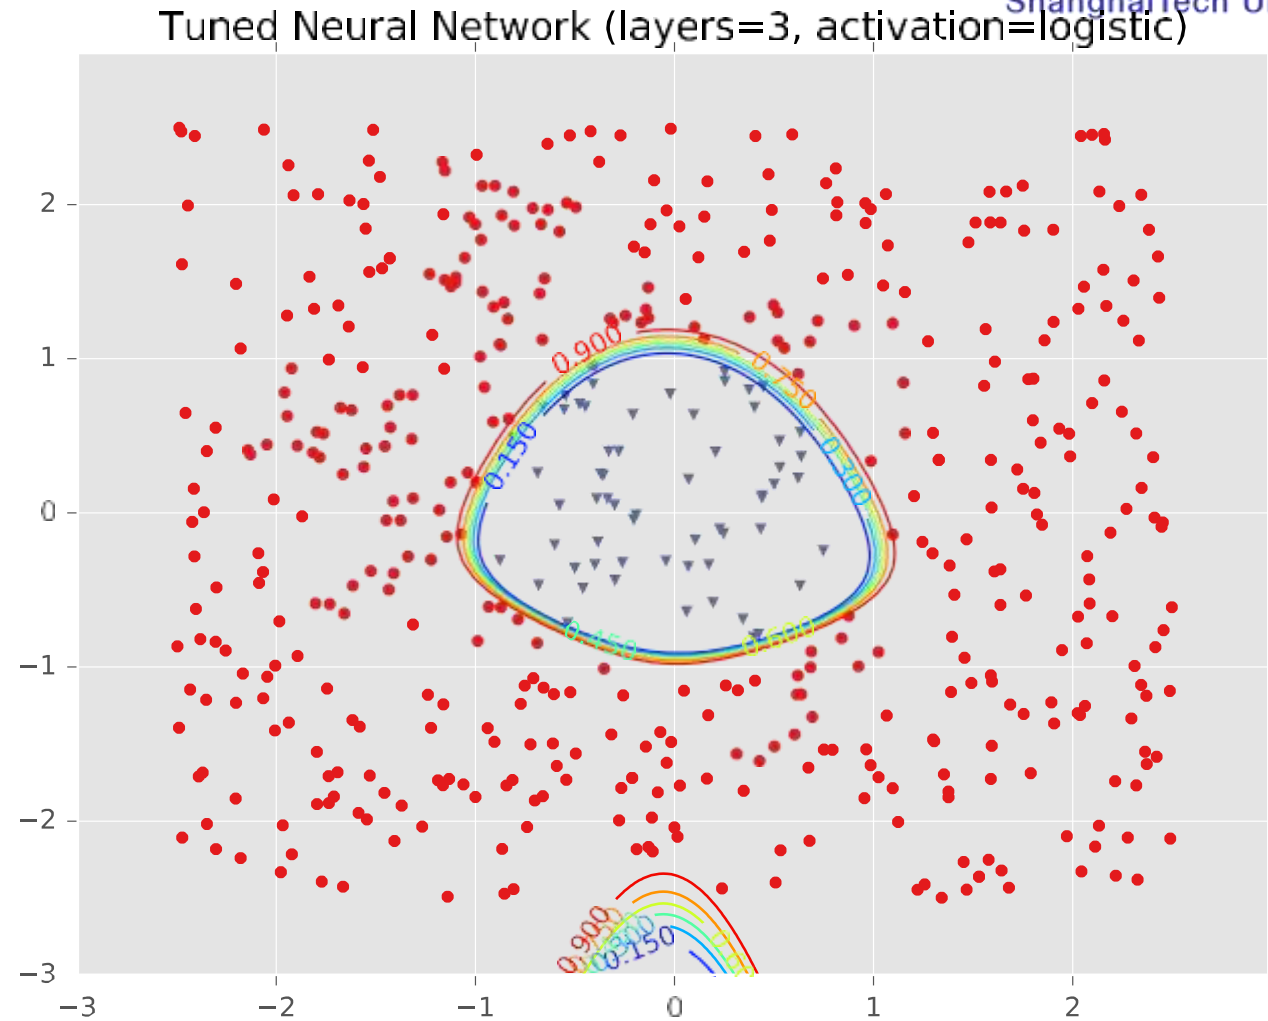


Figure courtesy of Matt Gormley

# Neural Network Decision Boundaries: Example 2



# Neural Network Decision Boundaries: Example 2



# Neural Network Decision Boundaries: Example 2

Tuned Neural Network (layers=3, activation=logistic)

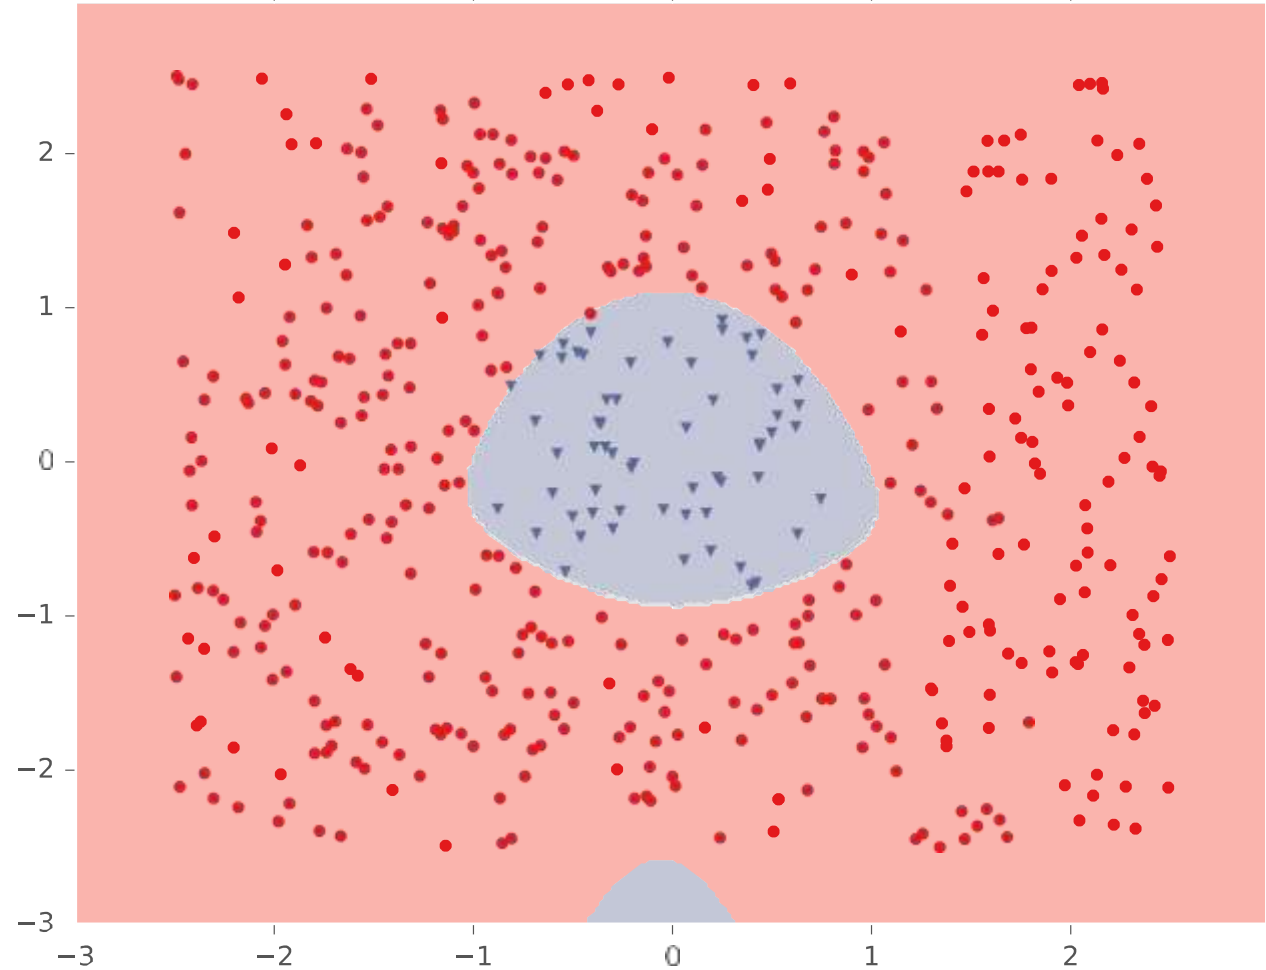
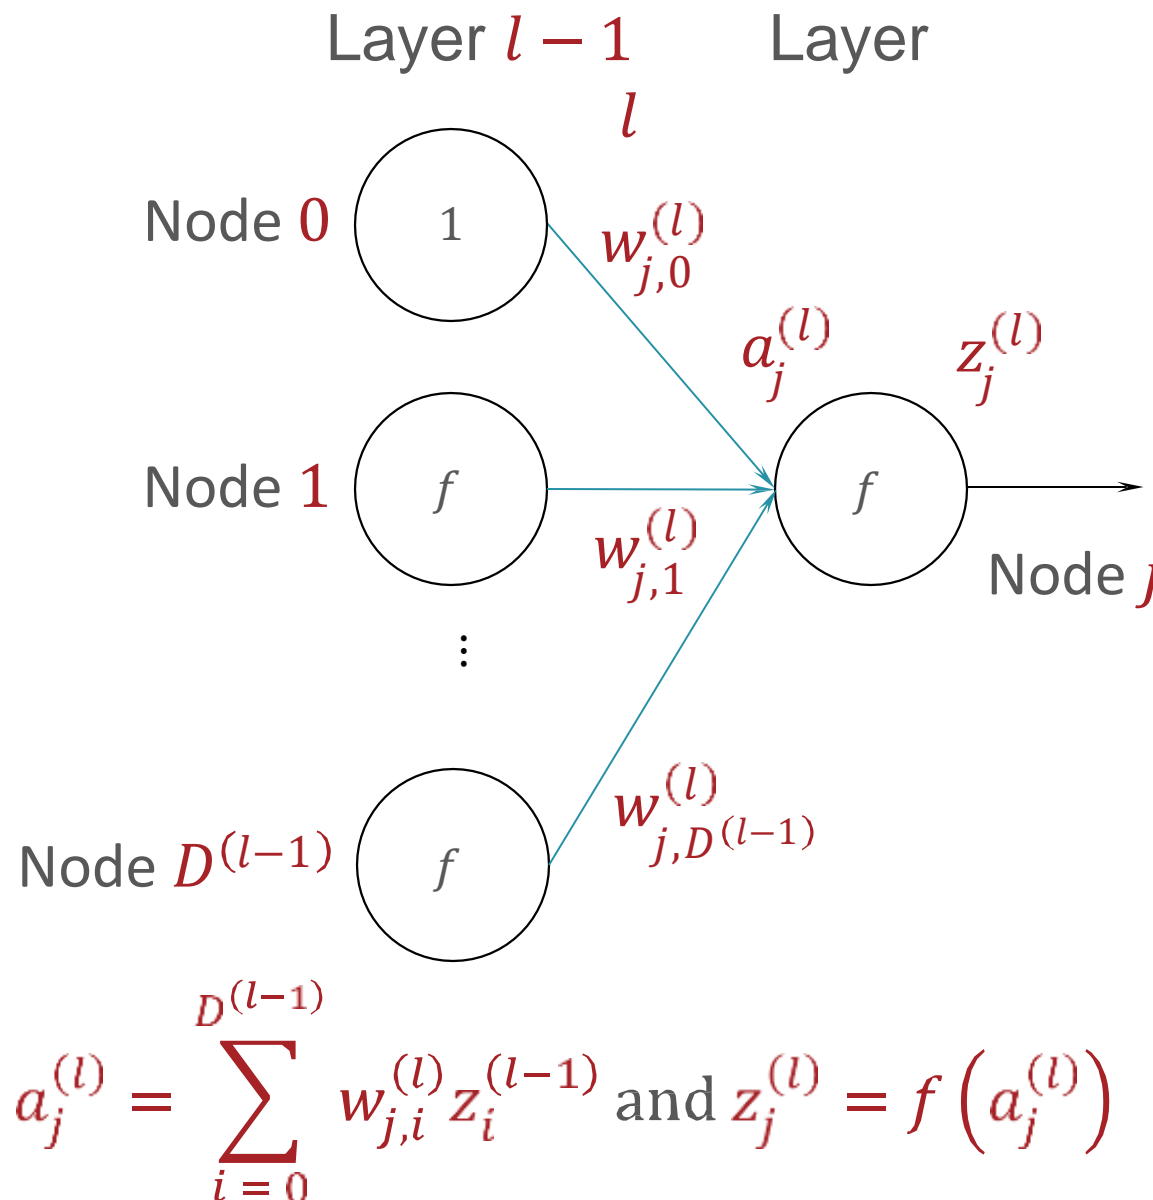


Figure courtesy of Matt Gormley

Every node has an incoming *signal* and outgoing *output*



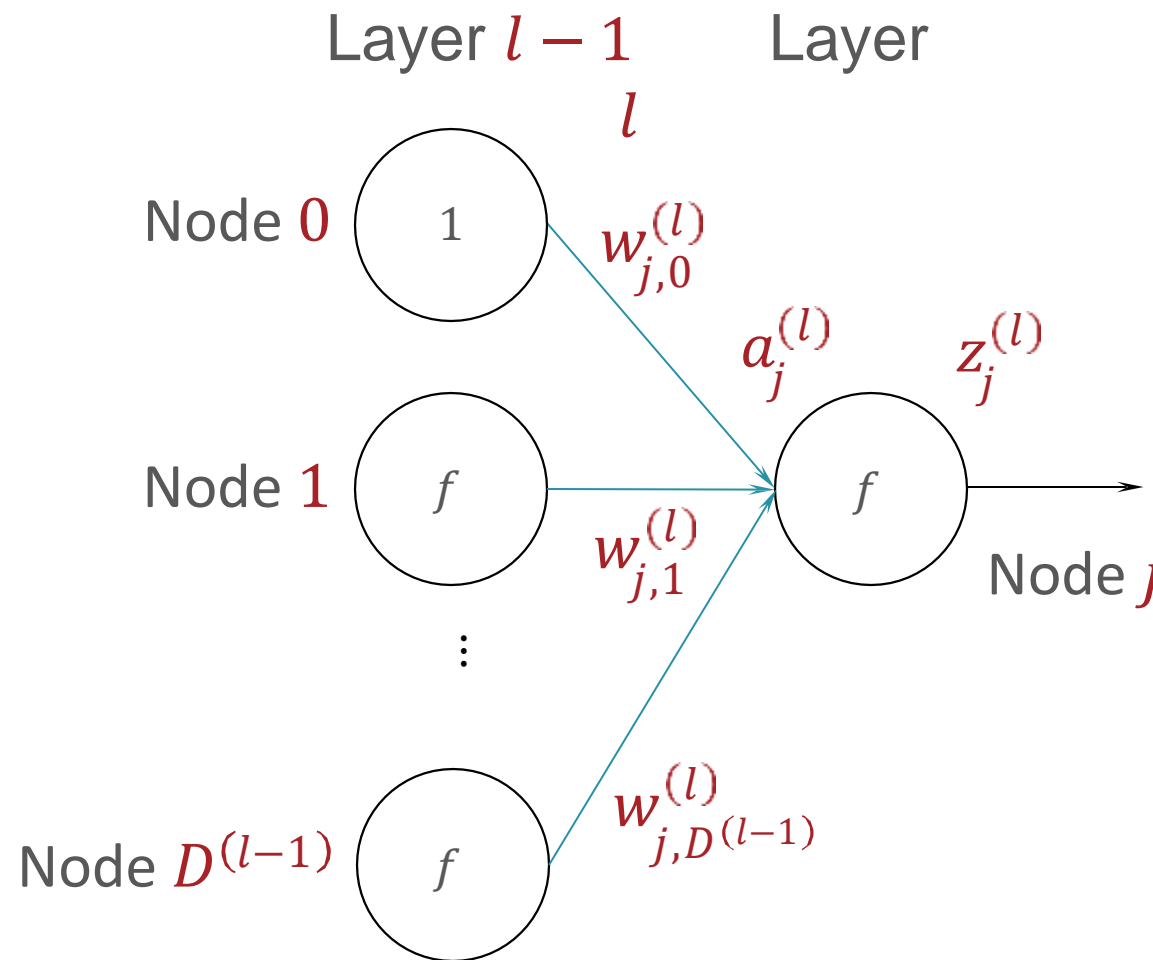
## Signal and Outputs



Every node has an incoming *signal* and outgoing *output*



## Signal and Outputs



$$\mathbf{a}^{(l)} = W^{(l)} \mathbf{z}^{(l-1)} \text{ and } \mathbf{z}^{(l)} = [1, f(\mathbf{z}^{(l)})]^T$$

## Forward Propagation for Making Predictions

- Input: weights  $W^{(1)}, \dots, W^{(L)}$  and a query data point  $\mathbf{x}$
- Initialize  $\mathbf{z}^{(0)} = [1, \mathbf{x}]^T$
- For  $l = 1, \dots, L$ 
  - $\mathbf{a}^{(l)} = W^{(l)} \mathbf{z}^{(l-1)}$
  - $\mathbf{z}^{(l)} = [1, f(\mathbf{a}^{(l)})]^T$
- Output:  $h_{W^{(1)}, \dots, W^{(L)}}(\mathbf{x}) = \mathbf{z}^{(L)}$



# Gradient Descent for Learning



- Input:  $\mathcal{D} = \{(\mathbf{x}^{(n)}, y^{(n)})\}_{n=1}^N, \eta^{(0)}$
- Initialize all weights  $W_{(0)}^{(1)}, \dots, W_{(0)}^{(L)}$  to small, random numbers and set  $t = 0$  (???)
- While TERMINATION CRITERION is not satisfied (???)
  - For  $l = 1, \dots, L$ 
    - Compute  $G^{(l)} = \nabla_{W^{(l)}} \ell_{\mathcal{D}}(W_{(t)}^{(1)}, \dots, W_{(t)}^{(L)})$  (???)
    - Update  $W^{(l)}$ :  $W_{(t+1)}^{(l)} = W_{(t)}^{(l)} - \eta_0 G^{(l)}$
  - Increment  $t$ :  $t = t + 1$
- Output:  $W_{(t)}^{(1)}, \dots, W_{(t)}^{(L)}$

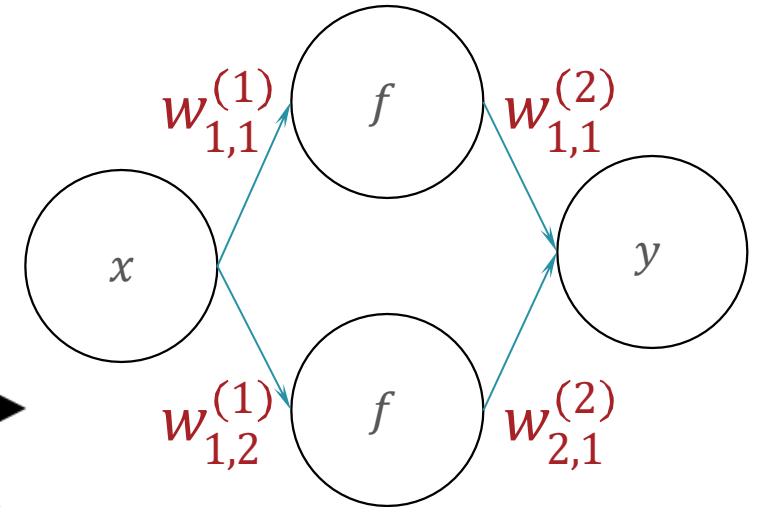
## Poll Question 2

- Suppose you are training a two-layer (one-hidden layer) neural network with sigmoid activations for binary classification.



- True or False: There is a unique set of parameters that maximizes the likelihood of the dataset above.

A. TOXIC      B. True      C. False



# Neural Network Learning Objectives

You should be able to...

1. Explain the biological motivations for a neural network
2. Combine simpler models (e.g. linear regression, binary logistic regression, multinomial logistic regression) as components to build up feed-forward neural network architectures
3. Explain the reasons why a neural network can model nonlinear decision boundaries for classification
4. Compare and contrast feature engineering with learning features
5. Identify (some of) the options available when designing the architecture of a neural network
6. Implement a feed-forward neural network