

CS182 Introduction to Machine Learning

Recitation 5

2025.3.26

Outline

- review: probability & statistics
- Logistic Regression
- Metrics for binary classification

Review(Preview): Probability & Statistics

- 什么用到概率论与数理统计：
 - 概率论为机器学习提供了问题的假设
 - 回归和分类问题都可以描述为一个估计问题
 - 数据的分布往往服从正态分布
- 用哪些知识：
 - 常用的概率公式（条件概率、全概率、贝叶斯）
 - 常用的分布和他们的特殊性质（正态、泊松、两点、二项、均匀）
 - 常用的统计量（均值、方差、协方差）和他们的无偏估计

Probability

- $P(A, B) = P(A|B)P(B) = P(B|A)P(A)$
- Bayes' Rule

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

$P(A)$: prior probability of A

$P(B)$: Evidence

$P(A|B)$: posterior probability of A given B

$P(B|A)$: likelihood (info of data given prior cognition)

Probability

- LOTP

$$P(A) = \sum_{i=1}^n P(A|B_i)P(B_i)$$

$\{B_1, B_2, \dots, B_n\}$ is a partition of the sample space. 交集为空集, 并集为全集

Conditional Probability

条件概率也是概率

- Bayes' Rule

$$P(A|B, C) = \frac{P(B|A, C)P(A|C)}{P(B|C)}$$

- LOTP

$$P(A|C) = \sum_{i=1}^n P(A|B_i, C)P(B_i|C)$$

Let $\hat{P}(\cdot) = P(\cdot|C)$

Independent & Exclusive Events 独立与互斥

- Independent:

$$P(A, B) = P(A)P(B)$$

$$P(A|B) = P(A)$$

$$P(B|A) = P(B)$$

- Exclusive:

$$P(A \cup B) = P(A) + P(B)$$

$$\text{i.e. } P(A \cap B) = 0$$

Expectation

$$\mathbb{E}(X) = \int x f(x) dx$$

Linearity

- $\mathbb{E}(X + Y) = \mathbb{E}(X) + \mathbb{E}(Y)$
- $\mathbb{E}(aX) = a\mathbb{E}(X)$
- $\nabla \mathbb{E}(f) = \mathbb{E}(\nabla f)$
- $X \perp\!\!\!\perp Y \Rightarrow \mathbb{E}(XY) = \mathbb{E}(X)\mathbb{E}(Y)$
LOTUS(Law of the Unconscious Statistician)
- $\mathbb{E}(g(X)) = \int g(x)f(x)dx$

条件期望也是期望

Variance

- $\text{Var}(X) = \mathbb{E}[(X - \mathbb{E}(X))^2]$
- $\text{Var}(X) = \mathbb{E}(X^2) - \mathbb{E}(X)^2$

If X and Y are independent, then $\text{Var}(X + Y) = \text{Var}(X) + \text{Var}(Y)$

方差没有线性性

- $\text{Var}(X + Y) = \text{Var}(X) + \text{Var}(Y) + 2\text{Cov}(X, Y)$
- $\text{Var}(cX) = c^2\text{Var}(X)$

Covariance

- $\text{Cov}(X, Y) = \mathbb{E}[(X - \mathbb{E}(X))(Y - \mathbb{E}(Y))]$
- $\text{Cov}(X, Y) = \mathbb{E}(XY) - \mathbb{E}(X)\mathbb{E}(Y)$
- $\text{Var}(X + Y) = \text{Var}(X) + \text{Var}(Y) + 2\text{Cov}(X, Y)$

Var($X + Y$) = Var(X) + Var(Y) if X and Y are independent

Correlation 相关性

- $\rho = \text{Corr}(X, Y) = \frac{\text{Cov}(X, Y)}{\sqrt{\text{Var}(X)\text{Var}(Y)}} \in [-1, 1]$

只能刻画**线性**相关性, 且正负相关程度相同(正负相关).

X, Y 独立, 则 $\rho_{X,Y} = 0$. 但是 $\rho_{X,Y} = 0$ **不一定**独立.

e.g. $Y = X^2, X \sim N(0, 1) \Rightarrow \mathbb{E}(X) = 0, \text{Var}(X) = 1, \mathbb{E}(Y) = \mathbb{E}(X^2) = 1$

$$\text{Cov}(X, Y) = \mathbb{E}(XY) - \mathbb{E}(X)\mathbb{E}(Y) = 0$$

*Gaussian*分布独立 \Leftrightarrow 不相关.

Iterated Expectation(Adam's law)

For any r.v. X , Y and Z :

$$\mathbb{E}_X [\mathbb{E}_{Y|X}(Y|X)] = \mathbb{E}_Y(Y)$$

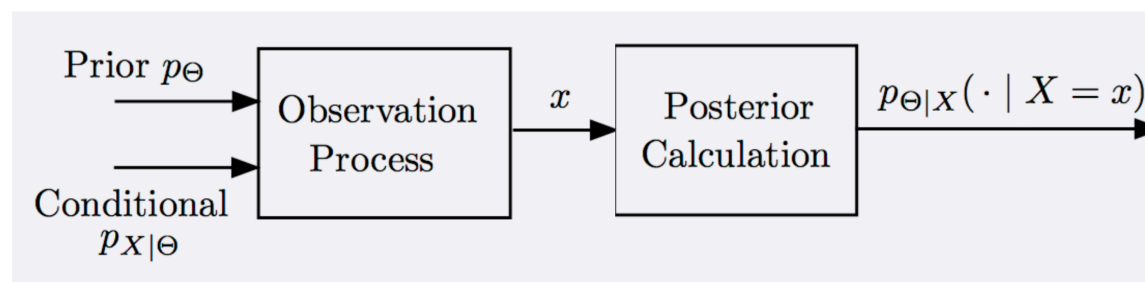
$$\mathbb{E}_X [\mathbb{E}_{Y|X}(Y|X, Z)|Z] = \mathbb{E}_Y(Y|Z)$$

Law of Total Variance(Eve's law)

$$\text{Var}(Y) = \mathbb{E}_X [\text{Var}(Y|X)] + \text{Var}(\mathbb{E}_{Y|X}(Y|X))$$

MAP: maximum a posteriori 最大后验分布

Bayesian 统计学派




贝叶斯学派把一切变量看作随机变量, 利用过去的知识和抽样数据,将概率解释为信念度 (degree of belief), 不需要大量的实验

Θ 是个随机变量, 根据样本的具体情况来估计参数, 使样本发生发可能性最大(与时俱进,不断更新)

$$\hat{\theta} = \arg \max_{\theta} p(\theta | \mathcal{D}) \propto P(\mathcal{D} | \theta) P_{\Theta}(\theta)$$

MLE: maximum a likelihood 最大似然估计

Frequentist 统计学派



频率学派把未知参数看作普通变量(固定值), 把样本看作随机变量, 仅仅利用抽样数据, 频率论方法通过大量独立实验将概率解释为统计均值(大数定律)

$$\hat{\theta} = \arg \max_{\theta} p(\mathcal{D}; \theta) \text{ or } \hat{\theta} = \arg \max_{\theta} p(\mathcal{D}|\theta)$$

$$\stackrel{i.i.d.}{=} \arg \max_{\theta} \prod_{i=1}^n p(y_i|\theta)$$

$$= \arg \max_{\theta} \sum_{i=1}^n \log p(y_i|\theta)$$

Conjugate Prior 共轭先验

Sample Space	Sampling Dist.	Conjugate Prior	Posterior
$\mathcal{X} = \{0, 1\}$	Bernoulli(θ)	Beta(α, β)	Beta($\alpha + n\bar{X}, \beta + n(1 - \bar{X})$)
$\mathcal{X} = \mathbb{Z}_+$	Poisson(λ)	Gamma(α, β)	Gamma($\alpha + n\bar{X}, \beta + n$)
$\mathcal{X} = \mathbb{Z}_{++}$	Geometric(θ)	Gamma(α, β)	Gamma($\alpha + n, \beta + n\bar{X}$)
$\mathcal{X} = \mathbb{H}_K$	Multinomial(θ)	Dirichlet(α)	Dirichlet($\alpha + n\bar{X}$)


Conjugate Prior 共轭先验

Sampling Dist.	Conjugate Prior	Posterior
Uniform(θ)	Pareto(ν_0, k)	Pareto ($\max\{\nu_0, X_{(n)}\}, n + k$)
Exponential(θ)	Gamma(α, β)	Gamma($\alpha + n, \beta + n\bar{X}$)
$N(\mu, \sigma^2)$, known σ^2	$N(\mu_0, \sigma_0^2)$	$N\left(\left(\frac{1}{\sigma_0^2} + \frac{n}{\sigma^2}\right)^{-1} \left(\frac{\mu_0}{\sigma_0^2} + \frac{n\bar{X}}{\sigma^2}\right), \left(\frac{1}{\sigma_0^2} + \frac{n}{\sigma^2}\right)^{-1}\right)$
$N(\mu, \sigma^2)$, known μ	InvGamma(α, β)	InvGamma $\left(\alpha + \frac{n}{2}, \beta + \frac{n}{2}\overline{(X - \mu)^2}\right)$
$N(\mu, \sigma^2)$, known μ	ScaledInv- $\chi^2(\nu_0, \sigma_0^2)$	ScaledInv- $\chi^2\left(\nu_0 + n, \frac{\nu_0\sigma_0^2}{\nu_0 + n} + \frac{n\overline{(X - \mu)^2}}{\nu_0 + n}\right)$
$N(\boldsymbol{\mu}, \boldsymbol{\Sigma})$, known $\boldsymbol{\Sigma}$	$N(\boldsymbol{\mu}_0, \boldsymbol{\Sigma}_0)$	$N\left(\mathbf{K}\left(\boldsymbol{\Sigma}_0^{-1}\boldsymbol{\mu}_0 + n\boldsymbol{\Sigma}^{-1}\bar{X}\right), \mathbf{K}\right), \mathbf{K} = (\boldsymbol{\Sigma}_0^{-1} + n\boldsymbol{\Sigma}^{-1})^{-1}$
$N(\boldsymbol{\mu}, \boldsymbol{\Sigma})$, known $\boldsymbol{\mu}$	InvWishart(ν_0, \mathbf{S}_0)	InvWishart($\nu_0 + n, \mathbf{S}_0 + n\bar{\mathbf{S}}$), $\bar{\mathbf{S}}$ sample covariance

Linear Least Square Estimate (LLSE)

拟合函数为线性: $h(X) = a + bX$, 则 minimizes $\mathbb{E} \left[(Y - a - bX)^2 \right]$ 的函数为


$$L[Y|X] = \mathbb{E}(Y) + \frac{\text{Cov}(X, Y)}{\text{Var}(X)} (X - \mathbb{E}(X))$$



Minimum Mean Square Error (MMSE)

拟合函数为**任意函数** $h(X)$, 则 minimizes $\mathbb{E} \left[(Y - h(X))^2 \right]$ 的函数为

$$h(X) = \mathbb{E}(Y|X)$$



Concentration Inequalities

For any r.v. X , and any constant $a > 0$:

- Markov Inequality

$$P(X \geq a) \leq \frac{\mathbb{E}(X)}{a}$$

- Chebyshev Inequality

$$P(|X - \mathbb{E}(X)| \geq a) \leq \frac{\text{Var}(X)}{a^2}$$

- Chernoff's Inequality

$$P(X \geq a) \leq \frac{\mathbb{E}(e^{tX})}{e^{ta}}, \forall t > 0$$

Probability Inequalities


- Hoeffding Lemma: $\mathbb{E}(X) = 0, a \leq X \leq b, \forall \lambda > 0$

$$\mathbb{E}(e^{\lambda X}) \leq e^{\frac{1}{8}\lambda^2(b-a)^2}$$

- Hoeffding Bound: X_1, \dots, X_n independent, $\mathbb{E}(X_i) = \mu, a \leq X_i \leq b, \forall \epsilon > 0$


$$P\left(\left|\frac{1}{n} \sum_{i=1}^n X_i - \mu\right| \geq \epsilon\right) \leq 2 \exp\left(-\frac{2n\epsilon^2}{\sum_{i=1}^n (b_i - a_i)^2}\right)$$

Logistic Regression



<https://zhuanlan.zhihu.com/p/139122386>

Sigmoid function



$$p(y_i = 1|x_i; \theta) = \sigma(\theta^\top x_i) = \frac{1}{1 + e^{-\theta^\top x_i}}$$

$$p(y_i = 0|x_i; \theta) = 1 - \sigma(\theta^\top x_i) = \frac{e^{-\theta^\top x_i}}{1 + e^{-\theta^\top x_i}}$$

Logistic Regression

$$p(y_i = 1|x_i; \theta) = \sigma(\theta^\top x_i) = \frac{1}{1 + e^{-\theta^\top x_i}}$$

$$p(y_i = 0|x_i; \theta) = 1 - \sigma(\theta^\top x_i) = \frac{e^{-\theta^\top x_i}}{1 + e^{-\theta^\top x_i}}$$

$$\Rightarrow p(y_i|x_i; \theta) = [\sigma(\theta^\top x_i)]^{y_i} \cdot [1 - \sigma(\theta^\top x_i)]^{1-y_i}$$

$$\hat{\theta} = \arg \max_{\theta} \prod_{i=1}^n p(y_i|x_i; \theta)$$

$$= \arg \max_{\theta} \prod_{i=1}^n [\sigma(\theta^\top x_i)]^{y_i} \cdot [1 - \sigma(\theta^\top x_i)]^{1-y_i}$$

$$= \arg \max_{\theta} \sum_{i=1}^n y_i \log [\sigma(\theta^\top x_i)] + (1 - y_i) \log [1 - \sigma(\theta^\top x_i)]$$

Multi-class Logistic Regression

sigmoid -> softmax

$$p(y_i = k|x_i; W) = \frac{e^{w_k^\top x_i}}{\sum_{j=1}^K e^{w_j^\top x_i}}$$

$$\mathcal{L}(W) = \prod_{i=1}^n \prod_{k=1}^K p(y_i = k|x_i; W)^{y_{ik}}$$

Or commonly as

$$\mathcal{L}(W) = \sum_{i=1}^n \sum_{k=1}^K y_{ik} \log p(y_i = k|x_i; W)$$

其中 $y_{ik} = 1$ 表示第 i 个样本的真实标签是第 k 类, 否则 $y_{ik} = 0$

Metrics for binary classification

- Accuracy
- Precision
- Recall
- F1 Score
- ROC Curve
- ROC-AUC

Metrics for binary classification

	Actual		
		Positive	Negative
		True Positive (TP)	False Positive (FP)
Predicted	Positive		
	Negative		
		False Negative (FN)	True Negative (TN)

True Positive(TP): 真阳性

False Positive(FP): 假阳性(错误的判断为正样本)

True Negative(TN): 真阴性

False Negative(FN): 假阴性(错误的判断为负样本)

reference: <https://zhuanlan.zhihu.com/p/364253497>


Metrics for binary classification

	Actual		
		Positive	Negative
		True Positive (TP)	False Positive (FP)
Predicted	Positive		
	Negative		
		False Negative (FN)	True Negative (TN)

- Accuracy: 正确率 = $\frac{TP+TN}{TP+TN+FP+FN}$
- Precision: 精准度 = $\frac{TP}{TP+FP}$
- Recall: 召回率 = $\frac{TP}{TP+FN}$
- F1 Score = $\frac{2 \times \text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$

Metrics for binary classification

- ROC (Receiver Operating Characteristic)



ROC-AUC(area-under-curve)